

Artificial Intelligence

Search (Chapter 3)

Instructor: Qiang Yu


Rubik's Cube robot

- <https://www.youtube.com/watch?v=iBE46R-fD6M>


Search

- We have some actions that can change the **state** of the world
 - Change induced by an action is perfectly predictable
- Try to come up with a sequence of actions that will lead us to a **goal state**
 - May want to minimize number of actions
 - More generally, may want to minimize total cost of actions
- Do not need to execute actions in real life while searching for solution!
 - Everything perfectly predictable anyway

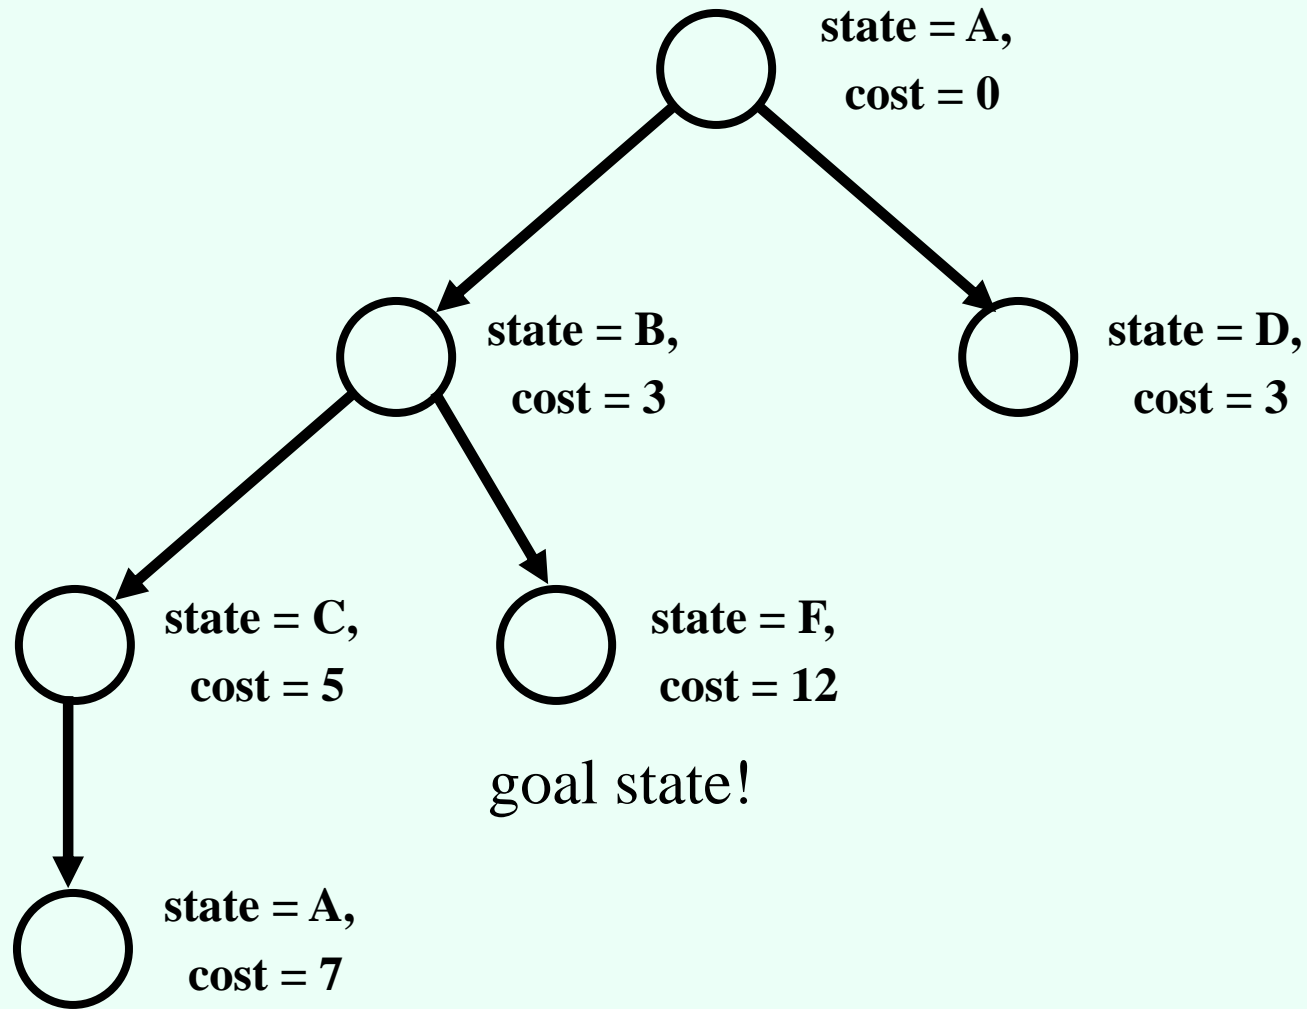
A simple example: traveling on a graph



Searching for a solution

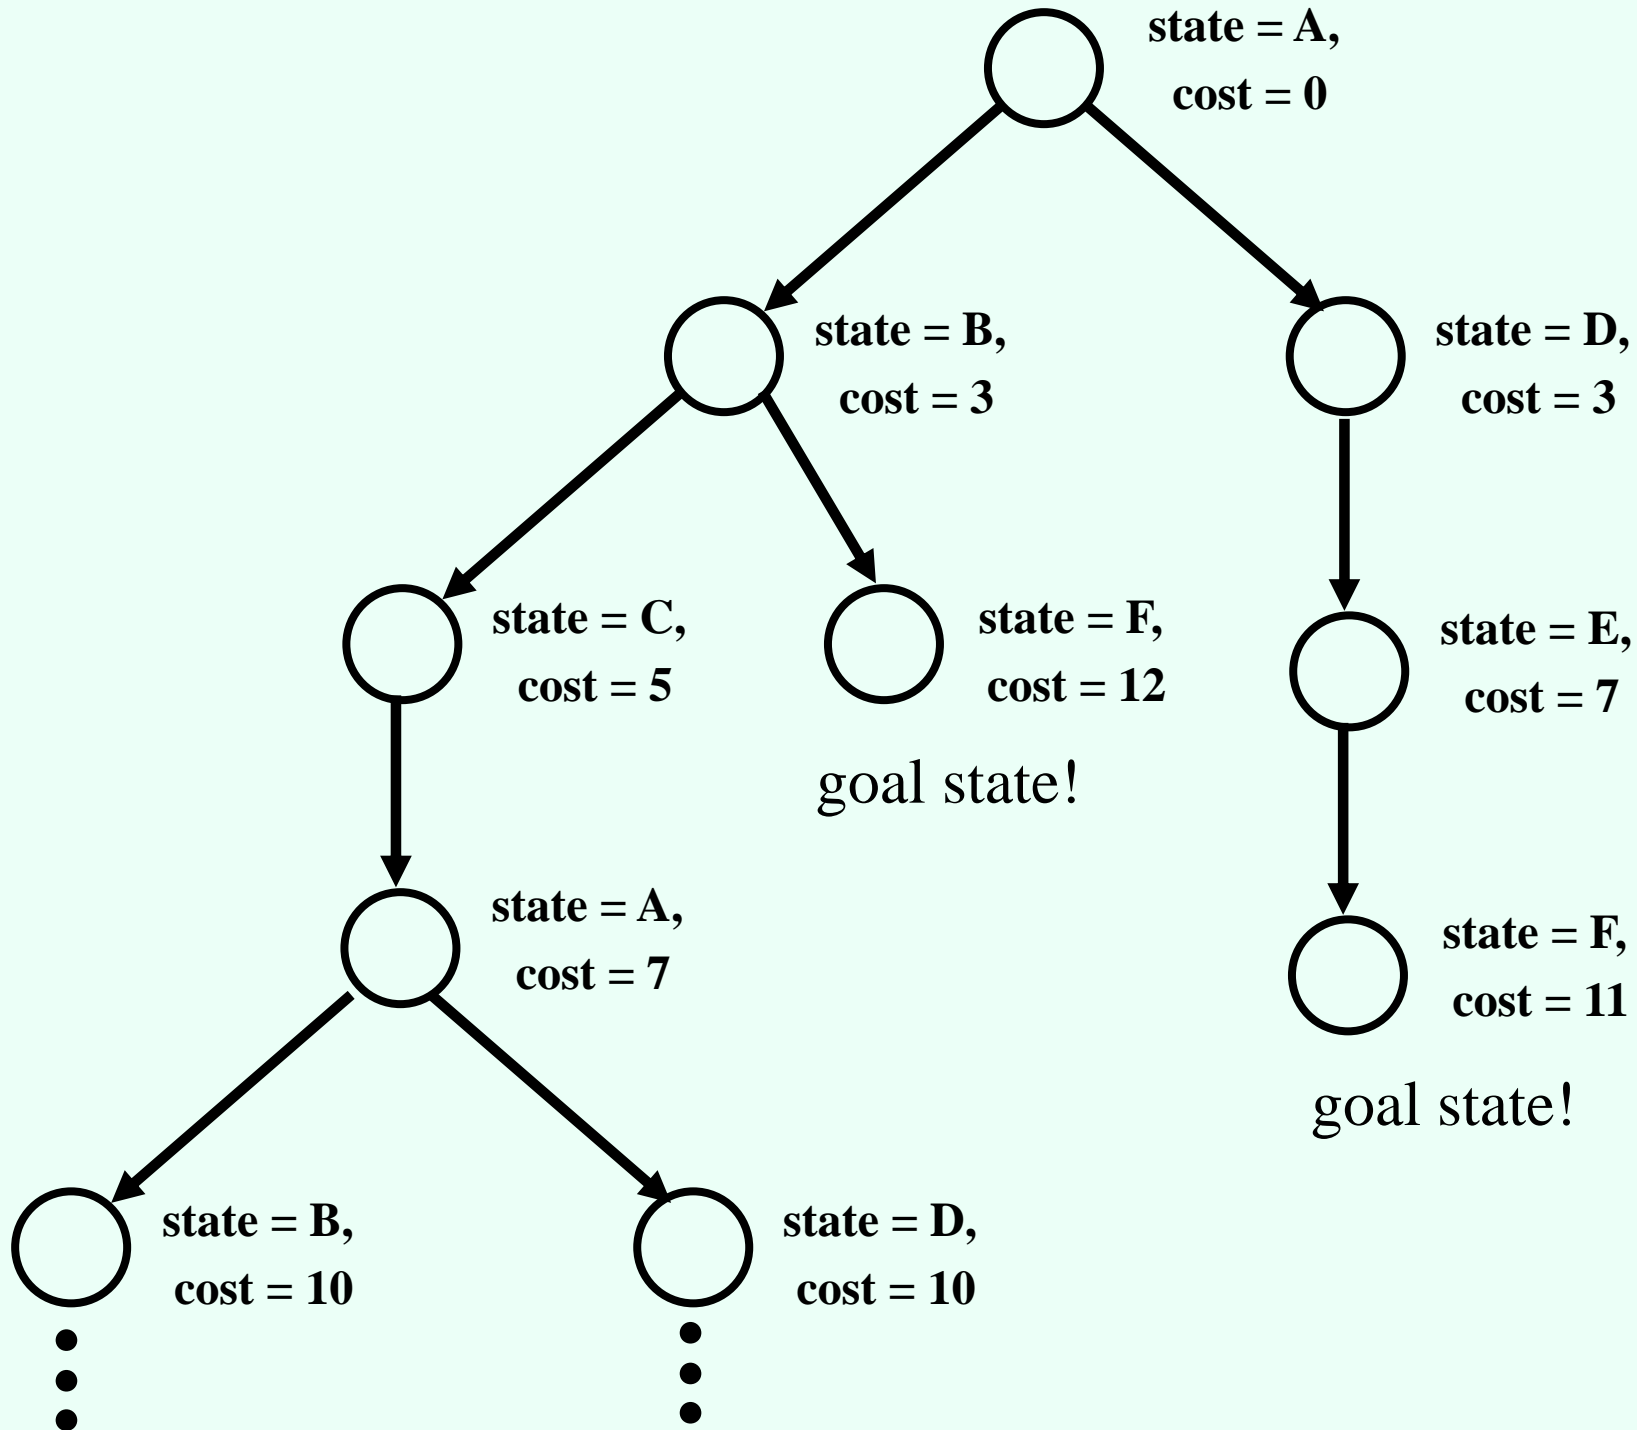


Search tree




search tree nodes and states are not the same thing!

Full search tree



Changing the goal:

want to visit all vertices on the graph



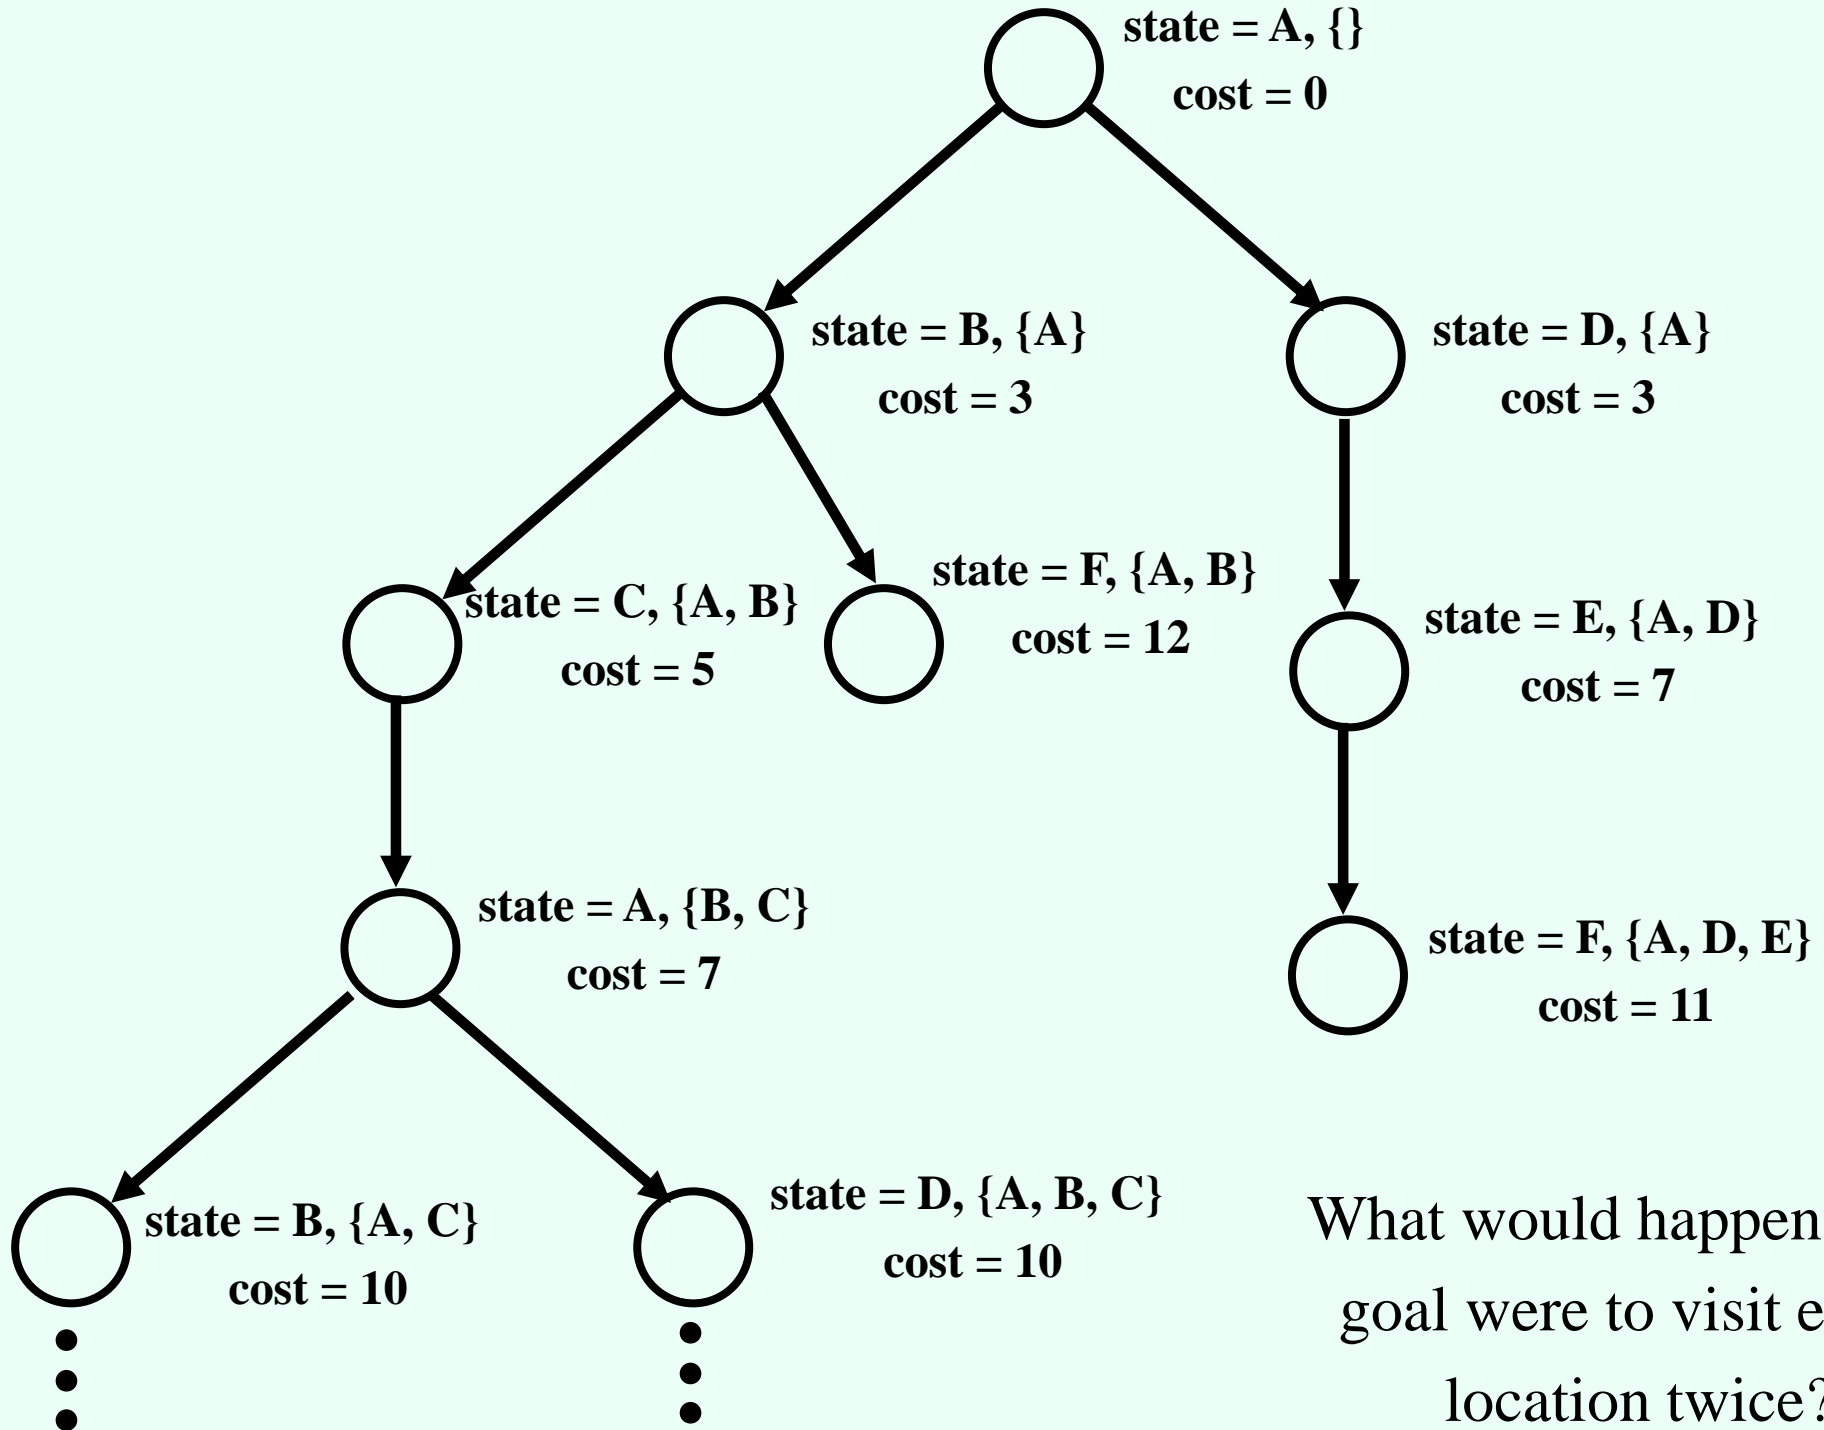
need a different definition of a state

“currently at A, also visited B, C already”

large number of states: $n * 2^{n-1}$

could turn these into a graph, but...

Full search tree



What would happen if the
goal were to visit every
location twice?

Key concepts in search

- Set of **states** that we can be in
 - Including an **initial state**...
 - ... and **goal states** (equivalently, a **goal test**)
- For every state, a set of **actions** that we can take
 - Each action results in a new state
 - Typically defined by **successor function**
 - Given a state, produces all states that can be reached from it
- **Cost function** that determines the cost of each action (or **path** = sequence of actions)
- **Solution**: path from initial state to a goal state
 - **Optimal solution**: solution with minimal cost


8-puzzle

1		2
4	5	3
7	8	6

1	2	3
4	5	6
7	8	

goal state

8-puzzle




Generic search algorithm

- **Frontier** = set of nodes **generated** but not **expanded**
- $\text{frontier} := \{\text{node with initial state}\}$
- loop:
 - if frontier empty, declare failure
 - choose and remove a node v from frontier
 - check if v 's state s is a goal state; if so, declare success
 - if not, expand v , insert resulting nodes into frontier
- Key question in search: Which of the generated nodes do we expand next?

Uninformed search

- Given a state, we only know whether it is a goal state or not
- Cannot say one nongoal state looks better than another nongoal state
- Can only traverse state space blindly in hope of somehow hitting a goal state at some point
 - Also called **blind search**
 - Blind does **not** imply unsystematic!

Breadth-first search



Properties of breadth-first search


- Nodes are expanded in the same order in which they are generated
 - Frontier can be maintained as a First-In-First-Out (FIFO) queue
- BFS is **complete**: if a solution exists, one will be found
- BFS finds a **shallowest** solution
 - Not necessarily an optimal solution
- If every node has b successors (the **branching factor**), first solution is at depth d , then frontier size will be at least b^d at some point
 - This much space (and time) required ☹

Properties of breadth-first search

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

Figure 3.13 Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

Depth-first search



Implementing depth-first search

- Frontier can be maintained as a Last-In-First-Out (LIFO) queue (aka. a stack)
- Also easy to implement recursively:
- DFS(node)
 - If goal(node) return solution(node);
 - For each successor of node
 - Return DFS(successor) unless it is *failure*;
 - Return *failure*;

Properties of depth-first search

- Not complete (might cycle through nongoal states)
- If solution found, generally not optimal/shallowest
- If every node has b successors (the **branching factor**), and we search to at most depth m , frontier is at most b^m
 - Much better space requirement 😊
 - Actually, generally don't even need to store all of frontier
- Time: still need to look at every node
 - $b^m + b^{m-1} + \dots + 1$ (for $b > 1$, $O(b^m)$)
 - **Inevitable** for uninformed search methods...

Combining good properties of BFS and DFS

- **Limited depth DFS:** just like DFS, except never go deeper than some depth d
- **Iterative deepening DFS:**
 - Call limited depth DFS with depth 0;
 - If unsuccessful, call with depth 1;
 - If unsuccessful, call with depth 2;
 - Etc.
- Complete, finds shallowest solution
- Space requirements of DFS
- May seem wasteful timewise because replicating effort
 - Really not that wasteful because **almost all effort at deepest level**
 - $db + (d-1)b^2 + (d-2)b^3 + \dots + 1b^d$ is $O(b^d)$ for $b > 1$

Let's start thinking about cost


- BFS finds shallowest solution because always works on shallowest nodes first
- Similar idea: always work on the **lowest-cost node** first (**uniform-cost** search)
- Will find optimal solution (assuming costs increase by at least constant amount along path)
- Will often pursue lots of short steps first
- If optimal cost is C , and cost increases by at least L each step, we can go to depth C/L
- Similar memory problems as BFS

Searching backwards from the goal

- Sometimes can search backwards from the goal
 - Maze puzzles
 - Eights puzzle
 - Reaching location F
 - What about the goal of “having visited all locations”?
- Need to be able to compute predecessors instead of successors
- What's the point?

Predecessor branching factor can be smaller than successor branching factor

- Stacking blocks:
 - only action is to add something to the stack



We'll see more of this...

Bidirectional search

- Even better: search from both the start and the goal, in parallel!





image from cs-alb-pc3.massey.ac.nz/notes/59302/fig03.17.gif

- If the shallowest solution has depth d and branching factor is b on both sides, requires only $O(b^{d/2})$ nodes to be explored!


Making bidirectional search work

- Need to be able to figure out whether the frontier intersect
 - Need to keep at least one frontier in memory...
- Other than that, can do various kinds of search on either tree, and get the corresponding optimality etc. guarantees
- Not possible (feasible) if backwards search not possible (feasible)
 - Hard to compute predecessors
 - High predecessor branching factor
 - Too many goal states

Repeated states



cycles



exponentially large search trees (try it!)

- Repeated states can cause incompleteness or enormous runtimes
- Can maintain list of previously visited states to avoid this
 - If new path to the same state has greater cost, don't pursue it further
 - Leads to time/space tradeoff
- “Algorithms that forget their history are doomed to repeat it” [Russell and Norvig]