

Artificial Intelligence

Search - more (Chapter 3)


Instructor: Qiang Yu

Informed search

- So far, have assumed that no nongoal state looks better than another
- Unrealistic
 - Even without knowing the road structure, some locations seem closer to the goal than others
 - Some states of the 8s puzzle seem closer to the goal than others
- Makes sense to expand closer-seeming nodes first

Heuristics

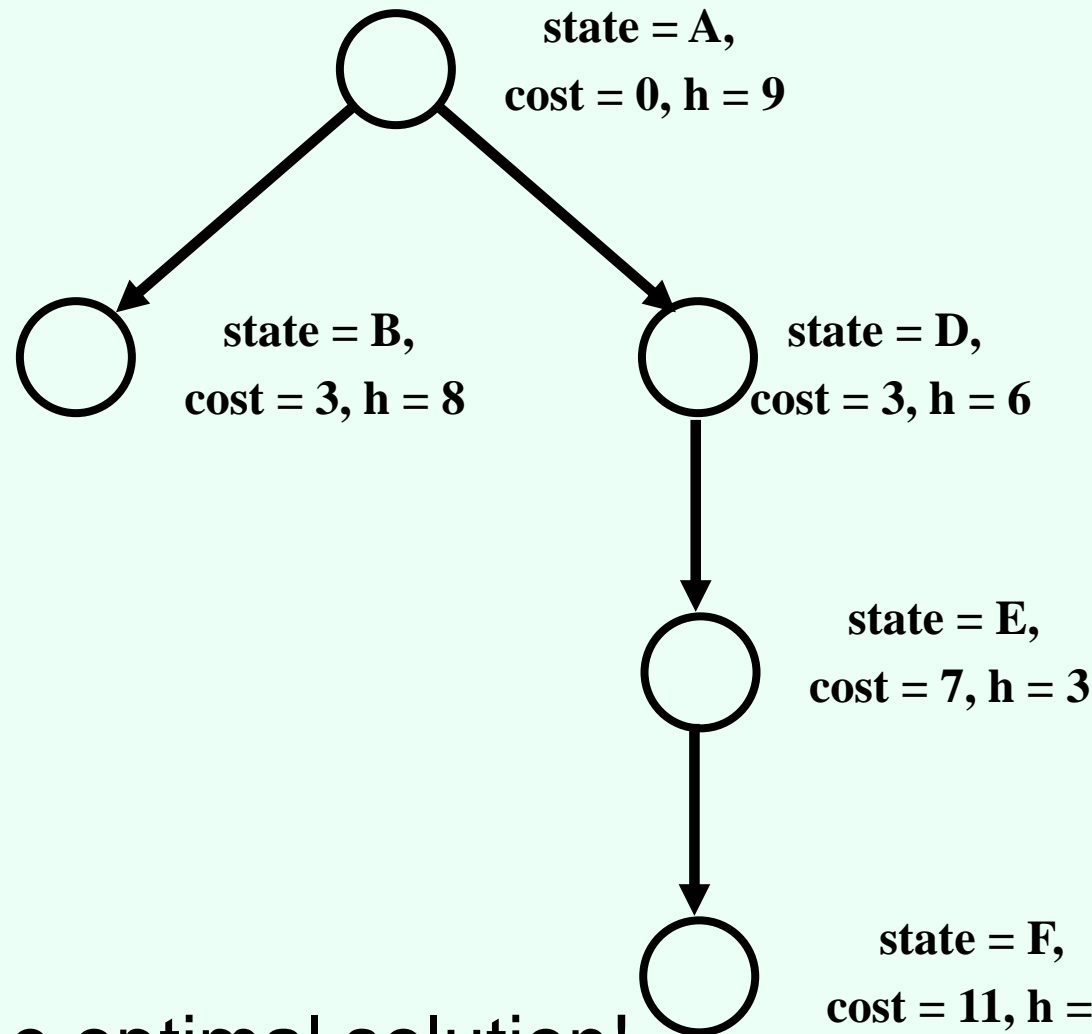
- Key notion: **heuristic function** $h(n)$ gives an estimate of the distance from n to the goal
 - $h(n)=0$ for goal nodes
- E.g. **straight-line distance** for traveling problem



- Say: $h(A) = 9$, $h(B) = 8$, $h(C) = 9$, $h(D) = 6$, $h(E) = 3$, $h(F) = 0$
- We're adding something new to the problem!
- Can use heuristic to decide which nodes to expand first


Greedy best-first search

- **Greedy best-first search:** expand nodes with lowest h values first



- Rapidly finds the optimal solution!
- Does it always?


A bad example for greedy



- Say: $h(A) = 9$, $h(B) = 5$, $h(D) = 6$, $h(E) = 3$, $h(F) = 0$
- Problem: greedy evaluates the promise of a node only by how far is left to go, does not take cost occurred already into account

A*

- Let $g(n)$ be cost incurred already on path to n
- Expand nodes with lowest $g(n) + h(n)$ first



- Say: $h(A) = 9$, $h(B) = 5$, $h(D) = 6$, $h(E) = 3$, $h(F) = 0$
- Note: if $h=0$ everywhere, then just uniform cost search

Admissibility

- A heuristic is **admissible** if it never overestimates the distance to the goal
 - If n is the optimal solution reachable from n' , then $g(n) \geq g(n') + h(n')$
- Straight-line distance is admissible: can't hope for anything better than a straight road to the goal
- Admissible heuristic means that A^* is always optimistic

Consistency

- A heuristic is **consistent** if the following holds: if one step takes us from n' to n , then $h(n') \leq h(n) + \text{cost of step from } n' \text{ to } n$
 - Similar to triangle inequality
 - Equivalently, $g(n') + h(n') \leq g(n) + h(n)$
- Implies admissibility

Iterative Deepening A*

- One big drawback of A* is the space requirement: similar problems as uniform cost search, BFS
- **Limited-cost depth-first A***: some cost cutoff c , any node with $g(n)+h(n) > c$ is not expanded, otherwise DFS

**More search:
When the path to the solution
doesn't matter**


(Chapter 4, 6)

Search where the path doesn't matter

- So far, looked at problems where the path was the solution
 - Traveling on a graph
 - Eights puzzle
- However, in many problems, we just want to find a goal state
 - Doesn't matter how we get there


Queens puzzle

- Place eight queens on a chessboard so that no two attack each other




Search formulation of the queens puzzle

- **Successors:** all valid ways of placing additional queen on the board; **goal:** eight queens placed




How big is this tree?
How many leaves?



Search formulation of the queens puzzle

- **Successors:** all valid ways of placing a queen in the next column; **goal:** eight queens placed



Search tree size?

What kind of search is best?

Constraint satisfaction problems (CSPs)


- Defined by:
 - A set of **variables** x_1, x_2, \dots, x_n
 - A **domain** D_i for each variable x_i
 - **Constraints** c_1, c_2, \dots, c_m
- A constraint is specified by
 - A subset (often, two) of the variables
 - All the allowable joint assignments to those variables
- Goal: find a **complete, consistent** assignment
- Queens problem: (other examples in next slides)
 - x_i in $\{1, \dots, 8\}$ indicates in which row in the i th column to place a queen
 - For example, constraint on x_1 and x_2 : $\{(1,3), (1,4), (1,5), (1,6), (1,7), (1,8), (2,4), (2,5), \dots, (3,1), (3,5), \dots \dots\}$

Meeting scheduling

- Meetings A, B, C, ... need to be scheduled on M, Tu, W, Th, F
- A and B cannot be scheduled on the same day
- B needs to be scheduled at least two days before C
- C cannot be scheduled on Th or F
- Etc.
- How do we model this as a CSP?

Graph coloring

- Fixed number of colors; no two adjacent nodes can share a color



Recap


- **Greedy best-first search**: expand nodes with lowest h values first
- **A* search**: $f=g+h$
- **Admissibility**: A heuristic is **admissible** if it never overestimates the distance to the goal.
- **Consistency**: if one step takes us from n' to n , then $h(n') \leq h(n) + \text{cost of step from } n' \text{ to } n$
- Constraint satisfaction problems (CSPs)

Generic approaches to solving CSPs

- State: some variables assigned, others not assigned
- Naïve successors definition: any way of assigning a value to an unassigned variable results in a successor
 - How many leaves do we get in the worst case?
- CSPs satisfy **commutativity**: order in which actions applied does not matter
- Better idea: only consider assignments for a single variable at a time
 - How many leaves?

Choice of variable to branch on is still flexible!

- Each of variables A , B , C takes values in $\{0,1\}$



A generic recursive search algorithm


(*assignment* is a partial assignment)

- **Search(*assignment*, *constraints*)**
- If *assignment* is complete, return it
- Choose an unassigned variable *x*
- For every value *v* in *x*'s domain, if setting *x* to *v* in *assignment* does not violate *constraints*:
 - Set *x* to *v* in *assignment*
 - *result* := Search(*assignment*, *constraints*)
 - If *result* != *failure* return *result*
 - Unassign *x* in *assignment*
- Return *failure*


Keeping track of remaining possible values

- For every variable, keep track of which values are still possible

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| | | | | Q | X | X | X |
| | | Q | | | X | X | X |
| Q | | | | | X | X | X |
| | | | | | | | X |
| | Q | | | | X | X | X |
| | | | | | X | X | |
| | | | | | | | X |
| | | | Q | | X | X | X |



| | | | | | | | |
|---|---|---|---|---|---|---|---|
| | | | | Q | X | X | |
| | | Q | | | X | X | |
| Q | | | | | X | X | |
| | | | | | X | | |
| | Q | | | | X | X | |
| | | | | | X | X | Q |
| | | | | | | X | |
| | | | Q | | X | X | |



| | | | | | | | |
|---|---|---|---|---|---|---|---|
| | | | | Q | | | |
| | | Q | | | | | |
| Q | | | | | | | |
| | | | | | | Q | |
| | Q | | | | | | |
| | | | | | | | Q |
| | | | | | Q | | |
| | | | Q | | | | |

only one possibility
for last column; might
as well fill in

now only one left for
other two columns

done!
(no real branching
needed!)

- General heuristic: branch on variable with fewest values remaining

Arc consistency

- Take two variables connected by a constraint
- Is it true that for **every** remaining value d of the first variable, there exists **some** value d' of the other variable so that the constraint is satisfied?
 - If so, we say the **arc from the first to the second variable is consistent**
 - If not, can remove the value d
- General concept: **constraint propagation**


| | | | | | | | |
|---|--|---|--|---|---|--|---|
| | | | | | | | |
| Q | | | | x | | | x |
| | | | | x | | | |
| | | | | | | | |
| | | Q | | x | | | x |
| | | | | x | | | |
| | | | | x | Q | | x |
| | | | | x | | | |

Is the arc from the fifth to the eighth column consistent?

What about the arc from the eighth to the fifth?


Tree-structured constraint graphs

- Suppose we only have pairwise constraints and the graph is a **tree** (or **forest** = multiple disjoint trees)



- Dynamic program for solving this (linear in #variables):
 - Starting from the leaves and going up, for each node x , compute all the values for x such that the subtree rooted at x can be solved
 - Equivalently: apply arc consistency from each parent to its children, starting from the bottom
 - If no domain becomes empty, once we reach the top, easy to fill in solution

Example: graph coloring with limited set of colors per node




- Stage 1: moving upward, cross out the values that cannot work with the subtree below that node
- Stage 2: if a value remains at the root, there is a solution: go downward to pick a solution

A different approach: optimization


- Let's say every way of placing 8 queens on a board, one per column, is **feasible**
- Now we introduce an **objective**: minimize the number of pairs of queens that attack each other
 - More generally, minimize the number of violated constraints
- Pure optimization

Local search: hill climbing


- Start with a complete state
- Move to successor with best (or at least better) objective value
 - Successor: move one queen within its column



4 attacking pairs




3 attacking pairs



2 attacking pairs

*no more
improvements*

- Local search can get stuck in a **local optimum**



Avoiding getting stuck with local search

- **Random restarts:** if your hill-climbing search fails (or returns a result that may not be optimal), restart at a random point in the search space
 - Not always easy to generate a random state
 - Will **eventually** succeed (why?)
- **Simulated annealing:**
 - Generate a random successor (possibly worse than current state)
 - Move to that successor with some probability that is sharply decreasing in the badness of the state
 - Also, over time, as the “temperature decreases,” probability of bad moves goes down