

## Chapter 18 solutions

### 18.1.1)

$r_1(A); r_1(B); w_1(B); r_1(C); w_1(C); r_1(D); w_1(D); r_1(E); w_1(E)$

### 18.1.2)

Think of the interleaving as having 10 positions. The 4 actions of the first transaction occupy any 4 of the 10, so the number of interleavings is (10 choose 4), or  $10*9*8*7/1*2*3*4 = 210$ .

### 18.2.1a)

Assume  $A=1, B=2$

#### **(T<sub>1</sub>, T<sub>2</sub>)**

##### **T<sub>1</sub>**

Read (A,t)	t=1
T=t+2	t=3
W(A,t)	A=3
R(B,t)	t=2
t=t*3	t=6
W(B,t)	B=6

##### **T<sub>2</sub>**

R(B,s)	s=6
s=s*2	s=12
W(B,s)	B=12
R(A,s)	s=3
s=s+3	s=6
W(A,s)	A=6

Final values after (T<sub>1</sub>, T<sub>2</sub>) : **A=6, B=12**

#### **(T<sub>2</sub>, T<sub>1</sub>)**

##### **T<sub>2</sub>**

R(B,s)	s=2
s=s*2	s=4
W(B,s)	B=4
R(A,s)	s=1
s=s+3	s=4
W(A,s)	A=4

##### **T<sub>1</sub>**

Read (A,t)	t=4
T=t+2	t=6
W(A,t)	A=6
R(B,t)	t=4
t=t*3	t=12
W(B,t)	B=12

Final values after (T<sub>2</sub>, T<sub>1</sub>) : **A=6, B=12**

**18.2.1b)**

Serializable schedule

<b>T<sub>1</sub></b> R(A,t) t=t*2 W(A,t)   R(B,t) t=t*3 W(B,t)	<b>T<sub>2</sub></b>   R(B,s) s=s*2 W(B,s)   R(A,s) s=s+3 W(A,s)
--	--

Non- serializable schedule

<b>T<sub>1</sub></b> R(A,t) t=t*2 W(A,t)  R(B,t)  t=t*3  W(B,t)	<b>T<sub>2</sub></b>   R(B,s)  s=s*2  W(B,s)  R(A,s) s=s+3 W(A,s)
--	--

The update of B by T<sub>2</sub> is lost. The final value of B is 6 (instead of 12 as in 18.2.1a)

**18.2.1c)**(T<sub>1</sub>,T<sub>2</sub>)(T<sub>2</sub>,T<sub>1</sub>)**18.2.1d)**

The key observation is that in order for an interleaving to be serializable, whichever transaction reads *A* first, must write it before the second reads *A*, and likewise for *B*.

There are four possibilities:

1. T<sub>1</sub> reads both *A* and *B* first: Then all the steps of T<sub>1</sub> must precede all the steps of T<sub>2</sub>, if the interleaving is serializable. There is 1 order of this type.
2. T<sub>2</sub> reads both *A* and *B* first: Similarly, there is 1 order of this type.
3. T<sub>1</sub> reads *B* first but T<sub>2</sub> reads *A* first. This situation is impossible.

4.  $T_1$  reads  $A$  first but  $T_2$  reads  $B$  first. Now, the first three steps of each transaction may interleave in any way, and the last three of each may interleave in any way, but the two groups of six actions must not interleave. The crucial observation is that the fourth step of each transaction (their second reads) must follow the third step of each transaction (their first writes), either to avoid reading  $A$  or  $B$  before it is written, or because actions of the same transaction cannot be reordered. The number of serializable orders of this type is  $(6 \text{ choose } 3) * (6 \text{ choose } 3) = 20 * 20 = 400$ .

The total number of serial orders is thus 402.

#### **18.2.2a)**

If a swap of two adjacent actions is legal, then the two actions can also be swapped back again, legally. Thus, instead of looking for schedules that are conflict equivalent to the given serial order, we can instead consider into what other schedules the serial order can be permuted by legal swaps. However, since in the serial order, the 4th and 5th steps are  $w_1(B); r_2(B)$ , these cannot be swapped. All other adjacent pairs are of the same transaction, and therefore they surely cannot be swapped. Thus, the only schedule conflict equivalent to the serial order  $(T_1, T_2)$  is that order itself; the answer is 1.

#### **18.2.2b)**

Only 1

$r_2(B)w_2(B)r_2(A)w_2(A)r_1(A)w_1(A)r_1(B)w_1(B)$

Fourth and fifth actions cannot be swapped since they conflict on  $A$  (atleast one is write)

#### **18.2.2c)**

Similar to 18.2.1 (d), for interleaving schedules to be serializable, a variable must be read and written by one transaction before another transaction reads it.

For serial schedule  $(T_1, T_2)$ , the four possibilities are:

1.  $T_1$  reads  $A$  &  $B$  first i.e. the one serial schedule  $(T_1, T_2)$ .
2.  $T_2$  reads  $A$  &  $B$  first. This is not possible for serial schedule  $(T_1, T_2)$ .
3.  $T_1$  reads  $B$  first while  $T_2$  reads  $A$  first. This is impossible since actions within a transaction cannot be reordered.
4.  $T_1$  reads  $A$  first while  $T_2$  reads  $B$  first. The first two steps of  $T_1$  can interleave in any way with first two steps of  $T_2$ . Similarly last two steps of  $T_1$  can interleave in any way with the last two steps of  $T_2$ . Thus number of serializable schedules are:  $(4 \text{ choose } 2) * (4 \text{ choose } 2) = 36$ .

Thus the total number of schedules =  $36 + 1 = 37$ .

#### **18.2.2d)**

Number of actions are different.

#### **18.2.3a)**

Conflict serializable schedules to  $(T_1, T_2)$

- 1)  $r_1(A)w_1(A)r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$

- 2)  $r_1(A)w_1(A)r_1(B)r_2(A)w_1(B)w_2(A)r_2(B)w_2(B)$  (2 swaps possible)
- 3)  $r_1(A)w_1(A)r_2(A)r_1(B)w_1(B)w_2(A)r_2(B)w_2(B)$  (first swap from 2)
- 4)  $r_1(A)w_1(A)r_2(A)r_1(B)w_2(A)w_1(B)r_2(B)w_2(B)$
- 5)  $r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$
- 6)  $r_1(A)w_1(A)r_1(B)r_2(A)w_2(A)w_1(B)r_2(B)w_2(B)$  (second swap from 2)

Similarly, there are six conflict-serializable schedules equivalent to serial order (T2,T1)

Hence total schedules =  $6 + 6 = 12$

### **18.2.3b)**

In addition to the twelve schedules above, interleaving using transaction behavior like 18.2.1 allows for two more possibilities:

1. T1 reads A first while T2 reads B first. Only such one schedule is possible:  
 $r_1(A);w_1(A);r_2(A);w_2(A);r_2(B);w_2(B);r_1(B);w_1(B);$

1. T1 reads B first while T2 reads A first. Only such one schedule is possible:  
 $r_2(A);w_2(A);r_1(A);w_1(A);r_1(B);w_1(B);r_2(B);w_2(B);$

Thus total number of serializable schedules =  $12 + 1 + 1 = 14$

### **18.2.4a)**

i)

The precedence graph is:

$T_3 \text{ -----} > T_2 \text{ -----} > T_1$

ii)

The schedule is thus conflict-serializable. The only conflict-equivalent serial schedule is (T<sub>3</sub>, T<sub>2</sub>, T<sub>1</sub>).

iii)

No. Assume, for example, that initially each of  $A$ ,  $B$ , and  $C$  are 10. T<sub>1</sub> sets  $A$  and  $C$  to 20, T<sub>2</sub> sets  $B$  to  $A+B+C$ , and T<sub>3</sub> prints  $B$ . Then if T<sub>3</sub> does not precede T<sub>2</sub>, it prints the wrong value, and if T<sub>2</sub> does not precede T<sub>1</sub>, it sets  $B$  to the wrong value.

### **18.2.4b)**

i)

$r_1(A)w_3(A) \quad T_1 < T_3$

$w_1(B)r_2(B) \quad T_1 < T_2$

$w_2(C)r_3(C) \quad T_2 < T_3$

Precedence graph is :

T1 -----> T2 -----> T3

ii)

Yes, schedule is conflict-serializable. Conflict-equivalent serial schedule is (T1,T2,T3).

iii)

No, same as part a iii)

#### **18.2.4c)**


i)

$w_3(A)r_1(A) \quad T_3 < T_1$

$w_1(B)r_2(B) \quad T_1 < T_2$

$w_2(C)r_3(C) \quad T_2 < T_3$

Precedence graph is:



ii)

No, schedule is not conflict-serializable

iii)

NA, since schedule has conflict.

#### **18.2.4d)**

i)


$r_1(A)r_2(A) \quad T_1 < T_2$

$w_1(B)w_2(B)r_1(B)r_2(B) \quad T_2 < T_1$

$w_2(C)$

w1(D)

Precedence graph is :



ii)

No, schedule is not conflict-serializable.

iii)

NA, since schedule has conflict.

### 18.2.4e)

i)

$r_1(A)r_2(A)r_3(A)w_1(A) \quad T_3 < T_1$

$r_1(B)r_2(B)r_4(B)w_2(B) \quad T_4 < T_2$

ii)

Yes, the schedule is conflict-serializable. The conflict-equivalent serial schedules are:

$T_3T_1T_4T_2$

$T_3T_4T_1T_2$

$T_3T_4T_2T_1$

$T_4T_3T_1T_2$

$T_4T_3T_2T_1$

$T_4T_2T_3T_1$

iii)

NA, since schedule has conflict.

### 18.2.5)

In order that the three conditions are satisfied, one approach can be

:T1 and T2 are not in conflict (better yet, not operating on same variables)

:Another transaction T3 is interleaved into the serial schedule (T1,T2) such that the actions of T3 that conflict with T1 precede it while the actions of T3 that conflict with T2 follow it.

Example:

T1:r1(A);W1(A);

T2:r2(B);w2(B);

T3:r3(A);w3(A);r3(B);w3(B);

Schedule S:

r3(A);w3(A);r1(A);w1(A);r2(B);w2(B);r3(B);w3(B);

The Precedence Graph of schedule S:

T2 -> T3 -> T1

Since T3 writes A before T1 reads A,  $T3 <_s T1$

Also, T2 writes B before T3 reads B,  $T2 <_s T3$

In schedule S, every action of T1 precedes T2. Hence condition i is satisfied.

There is no cycle in the Precedence Graph. Hence S is conflict serializable(condition ii)

Also from Precedence Graph, T2 precedes T1 for every serial schedule conflict-equivalent to S.

### **18.2.6)**

Consider the precedence graph of the serial schedule of transactions (T1,T2,T3,...) such that T1 has reads for all variables and all other transactions write two of the variables such that T1 is in conflict with all other transactions.

A cycle of length n can be found by reversing the conflicting actions of T1 and Tn. Since no transaction other than T1 and Tn operate on the variable in the conflicting actions, no smaller cycle exists.

### **18.3.1a)**

### **18.3.1b)**

In legal schedules locks on a particular element are only held by one transaction at a given time. Since locks are being acquired by T1 and T2 just before reads and being released just after writes; we can restate the legality requirement here as:

Between the read and write of an element by one transaction, No other transaction should act on the element. Thus below 32 transactions are illegal:

r2(B);w2(B);r2(A);r1(A);w2(A);w1(A);r1(B);w1(B);  
r2(B);w2(B);r2(A);r1(A);w1(A);w2(A);r1(B);w1(B);  
r2(B);w2(B);r2(A);r1(A);w1(A);r1(B);w2(A);w1(B);  
r2(B);w2(B);r2(A);r1(A);w1(A);r1(B);w1(B);w2(A);  
r2(B);w2(B);r1(A);r2(A);w2(A);w1(A);r1(B);w1(B);  
r2(B);w2(B);r1(A);r2(A);w1(A);w2(A);r1(B);w1(B);  
r2(B);w2(B);r1(A);r2(A);w1(A);r1(B);w2(A);w1(B);  
r2(B);w2(B);r1(A);r2(A);w1(A);r1(B);w1(B);w2(A);

r2(B);r1(A);w2(B);r2(A);w2(A);w1(A);r1(B);w1(B);  
r2(B);r1(A);w2(B);r2(A);w1(A);w2(A);r1(B);w1(B);  
r2(B);r1(A);w2(B);r2(A);w1(A);r1(B);w2(A);w1(B);  
r2(B);r1(A);w2(B);r2(A);w1(A);r1(B);w1(B);w2(A);

r2(B);r1(A);w1(A);r1(B);w2(B);r2(A);w2(A);w1(B);

r2(B);r1(A);w1(A);r1(B);w2(B);r2(A);w1(B);w2(A);  
r2(B);r1(A);w1(A);r1(B);w2(B);w1(B);r2(A);w2(A);  
r2(B);r1(A);w1(A);r1(B);w1(B);w2(B);r2(A);w2(A);  
r1(A);r2(B);w2(B);r2(A);w2(A);w1(A);r1(B);w1(B);  
r1(A);r2(B);w2(B);r2(A);w1(A);w2(A);r1(B);w1(B);  
r1(A);r2(B);w2(B);r2(A);w1(A);r1(B);w2(A);w1(B);  
r1(A);r2(B);w2(B);r2(A);w1(A);r1(B);w1(B);w2(A);

r1(A);r2(B);w1(A);r1(B);w2(B);r2(A);w2(A);w1(B);  
r1(A);r2(B);w1(A);r1(B);w2(B);r2(A);w1(B);w2(A);  
r1(A);r2(B);w1(A);r1(B);w2(B);w1(B);r2(A);w2(A);  
r1(A);r2(B);w1(A);r1(B);w1(B);w2(B);r2(A);w2(A);

r1(A);w1(A);r2(B);r1(B);w2(B);r2(A);w2(A);w1(B);  
r1(A);w1(A);r2(B);r1(B);w2(B);r2(A);w1(B);w2(A);  
r1(A);w1(A);r2(B);r1(B);w2(B);w1(B);r2(A);w2(A);  
r1(A);w1(A);r2(B);r1(B);w1(B);w2(B);r2(A);w2(A);  
r1(A);w1(A);r1(B);r2(B);w2(B);r2(A);w2(A);w1(B);  
r1(A);w1(A);r1(B);r2(B);w2(B);r2(A);w1(B);w2(A);  
r1(A);w1(A);r1(B);r2(B);w2(B);w1(B);r2(A);w2(A);  
r1(A);w1(A);r1(B);r2(B);w1(B);w2(B);r2(A);w2(A);

While below 38 transactions are legal:

r2(B);w2(B);r2(A);w2(A);r1(A);w1(A);r1(B);w1(B);

r2(B);w2(B);r1(A);w1(A);r2(A);w2(A);r1(B);w1(B);  
r2(B);w2(B);r1(A);w1(A);r2(A);r1(B);w2(A);w1(B);  
r2(B);w2(B);r1(A);w1(A);r2(A);r1(B);w1(B);w2(A);  
r2(B);w2(B);r1(A);w1(A);r1(B);r2(A);w2(A);w1(B);  
r2(B);w2(B);r1(A);w1(A);r1(B);r2(A);w1(B);w2(A);  
r2(B);w2(B);r1(A);w1(A);r1(B);w1(B);r2(A);w2(A);  
r2(B);r1(A);w2(B);w1(A);r2(A);w2(A);r1(B);w1(B);  
r2(B);r1(A);w2(B);w1(A);r2(A);r1(B);w2(A);w1(B);  
r2(B);r1(A);w2(B);w1(A);r2(A);r1(B);w1(B);w2(A);  
r2(B);r1(A);w2(B);w1(A);r1(B);r2(A);w2(A);w1(B);  
r2(B);r1(A);w2(B);w1(A);r1(B);w1(B);r2(A);w2(A);  
r2(B);r1(A);w1(A);w2(B);r2(A);w2(A);r1(B);w1(B);  
r2(B);r1(A);w1(A);w2(B);r2(A);r1(B);w2(A);w1(B);  
r2(B);r1(A);w1(A);w2(B);r2(A);r1(B);w1(B);w2(A);  
r2(B);r1(A);w1(A);w2(B);r1(B);r2(A);w2(A);w1(B);  
r2(B);r1(A);w1(A);w2(B);r1(B);r2(A);w1(B);w2(A);



r2(B);r1(A);w1(A);w2(B);r1(B);w1(B);r2(A);w2(A);

r1(A);r2(B);w2(B);w1(A);r2(A);w2(A);r1(B);w1(B);  
r1(A);r2(B);w2(B);w1(A);r2(A);r1(B);w2(A);w1(B);  
r1(A);r2(B);w2(B);w1(A);r2(A);r1(B);w1(B);w2(A);  
r1(A);r2(B);w2(B);w1(A);r1(B);r2(A);w2(A);w1(B);  
r1(A);r2(B);w2(B);w1(A);r1(B);r2(A);w1(B);w2(A);  
r1(A);r2(B);w2(B);w1(A);r1(B);w1(B);r2(A);w2(A);  
r1(A);r2(B);w1(A);w2(B);r2(A);w2(A);r1(B);w1(B);  
r1(A);r2(B);w1(A);w2(B);r2(A);r1(B);w2(A);w1(B);  
r1(A);r2(B);w1(A);w2(B);r2(A);r1(B);w1(B);w2(A);  
r1(A);r2(B);w1(A);w2(B);r1(B);r2(A);w1(B);w2(A);  
r1(A);r2(B);w1(A);w2(B);r1(B);w1(B);r2(A);w2(A);  
r1(A);w1(A);r2(B);w2(B);r2(A);w2(A);r1(B);w1(B);  
r1(A);w1(A);r2(B);w2(B);r2(A);r1(B);w2(A);w1(B);  
r1(A);w1(A);r2(B);w2(B);r2(A);r1(B);w1(B);w2(A);  
r1(A);w1(A);r2(B);w2(B);r1(B);r2(A);w2(A);w1(B);  
r1(A);w1(A);r2(B);w2(B);r1(B);r2(A);w1(B);w2(A);  
r1(A);w1(A);r2(B);w2(B);r1(B);w1(B);r2(A);w2(A);

r1(A);w1(A);r1(B);w1(B);r2(B);w2(B);r2(A);w2(A);

### **18.3.1c)**

According to semantics of transactions, four possibilities for serializable schedules are:

- (i) Serial order (T1,T2) = 1 transaction
- (ii) Serial order (T2,T1) = 1 transaction
- (iii) T1 reads B first while T2 reads A first. This is impossible.
- (iv) T2 reads A first while T1 reads B first. The first 2 actions of T1 can interleave with first two actions of T2. Similarly the last two actions of T1 can interleave with last two actions of T2.

But the two groups of 4 actions cannot interleave. Thus there are  $(4 \text{ choose } 2) * (4 \text{ choose } 2) = 36$ .

Hence total = 38 serializable schedules.

### **18.3.1d)**

Only two schedules are conflict serializable: the serial orders (T1,T2) and (T2,T1). Both T1 and T2 start with the read of one element and end with write of the other element. Thus no conflict-free swaps are possible in the serial orders.

### 18.3.1e)

No, there are none. Due to the semantics of the transactions, and the placement of lock and unlock operations (as explained in (b)), all legal schedules are serializable in this exercise.

### 18.3.2)

Suppose the schedule starts with  $T_1$  locking and reading  $A$ . If  $T_2$  locks  $B$  before  $T_1$  reaches its unlocking phase, then there is a deadlock, and the schedule cannot complete. Thus, if  $T_1$  performs an action first, it must perform *all* its actions before  $T_2$  performs any. Likewise, if  $T_2$  starts first, it must complete before  $T_1$  starts, or there is a deadlock. Thus, only the two serial schedules of these transactions are legal.

### 18.3.3a)

Note that this question refers to the solved Exercise 18.2.4(a). For the sequence of the latter exercise, the second request,  $r_2(A)$ , cannot be granted because of the lock held by  $T_1$ .  $T_2$  is thus delayed until after the 4th step of the sequence,  $w_1(A)$ , after which  $T_1$  releases its lock on  $A$ . No other delays are necessary, so the sequence of scheduled actions would be:  $r_1(A)$ ;  $r_3(B)$ ;  $w_1(A)$ ;  $r_2(A)$ ;  $r_2(C)$ ;  $r_2(B)$ ;  $w_2(B)$ ;  $w_1(C)$ .

### 18.3.3b)

$r_1(A)w_1(B)r_2(B)w_2(C)r_3(C)w_3(A)$

No delays necessary.

### 18.3.3c)

$w_3(A)r_1(A)w_1(B)r_2(B)w_2(C)r_3(C)$

No delays necessary.

### 18.3.3d)

$r_1(A)r_2(A)w_1(B)\underline{w_2(B)}r_1(B)r_2(B)w_2(C)w_1(D)$

$w_2(B)$  is delayed since  $w_1(B)$  is holding a lock to  $B$ .

Modified schedule:

$r_1(A)r_2(A)w_1(B)r_1(B)w_2(B)r_2(B)w_2(C)w_1(D)$

### 18.3.3.e)

$r_1(A)\underline{r_2(A)}r_1(B)r_2(B)\underline{r_3(A)}\underline{r_4(B)}w_1(A)w_2(B)$

$r_2(A)$  and  $r_3(A)$  are delayed since  $r_1(A)$  is holding a lock to  $A$ .

$r_4(B)$  is delayed since  $r_2(B)$  is holding a lock to  $B$ .

Modified schedule:

$r_1(A)r_1(B)r_2(B)w_1(A)r_2(A)r_3(A)w_2(B)r_4(B)$

### 18.3.4a)

The consistent sequences(i+ii) are interleavings of the sequences

1.  $l_1(A);r_1(A);u_1(A)$

2.  $l_1(B);w_1(B);u_1(B)$

The number of such interleaving are  $(6 \text{ choose } 3) = 20$ .

A sequence is not two-phase locked if the unlock of 1 occurs before lock of 2 or vice versa. Only 2 such interleavings are possible i.e. all of 1 followed by all of 2 or vice versa. Thus number of consistent but not two-phase orders = 2. (ii)

Thus the number of two-phase locked orders are =  $20 - 2 = 18$ . (i)

Next consider number of two-phased locked orders(i+iii). The two locks must occur before the two locks; there are four such permutations:

$l_2(A); l_2(B); u_2(A); u_2(B)$

$l_2(A); l_2(B); u_2(B); u_2(A)$

$l_2(B); l_2(A); u_2(A); u_2(B)$

$l_2(B); l_2(A); u_2(B); u_2(A)$

Action  $r_1(A)$  can be placed in any of 5 positions combining with above lock-unlock actions i.e. before, after, or in between above lock-unlock actions.

Similarly  $w_1(B)$  can be placed in any of 6 positions combining above lock-unlock- $r_1(A)$  i.e, before, after, in-between.

Thus total two phase lock orders are =  $4 * 5 * 6 = 120$ .

Thus number of inconsistent but two phase locked orders =  $120 - 18 = 102$ . (iii)

Total number of orders(i+ii+iii+iv) of 6 actions =  $6! = 720$

Thus number of inconsistent non-two phase locked orders =  $720 - 18 - 2 - 102 = 598$  (iv)

Thus the answers are

(i) 18

(ii) 2

(iii) 102

(iv) 598

### **18.3.4(b)**

The consistent sequences(i+ii) are interleavings of the sequences

1.  $l_2(A); r_2(A); w_2(A); u_2(A)$

2.  $l_2(B); w_2(B); u_2(B)$

The number of such interleaving are  $(7 \text{ choose } 3) = 35$ . Since in 1  $r_2(A)$  and  $w_2(A)$  can be interchanged, the possible consistent sequences are =  $35 * 2 = 70$ .

A sequence is not two-phase locked if the unlock of 1 occurs before lock of 2 or vice versa. Only 2 such interleavings are possible i.e. all of 1 followed by all of 2 or vice versa. Again since in 1  $r_2(A)$  and  $w_2(A)$  can be interchanged, total number of consistent but non-two-phase locked orders are =  $2 * 2 = 4$  (ii)

Thus the number of two-phase locked orders are =  $70 - 4 = 66$ . (i)

Next consider number of two-phased locked orders(i+iii). The two locks must occur before the two locks; there are four such permutations:

$l_2(A); l_2(B); u_2(A); u_2(B)$

$l_2(A); l_2(B); u_2(B); u_2(A)$

$l_2(B); l_2(A); u_2(A); u_2(B)$

$l_2(B); l_2(A); u_2(B); u_2(A)$

Action  $r_2(A)$  can be placed in any of 5 positions combining with above lock-unlock actions i.e. before, after, or in between above lock-unlock actions.

Similarly  $w_2(A)$  can be placed in any of 6 positions combining above lock-unlock- $r_2(A)$  i.e, before, after, in-between.

Similarly  $w_2(B)$  can be placed in any of 7 positions above. Thus total two phase lock orders are  $= 4 * 5 * 6 * 7 = 840$ .

Thus number of inconsistent but two phase locked orders  $= 840 - 66 = 774$ . (iii)

Total number of orders(i+ii+iii+iv) of 7 actions  $= 7! = 5040$

Thus number of inconsistent non-two phase locked orders  $= 5040 - 774 - 66 - 4 = 4196$  (iv)

Thus the answers are

(i) 66

(ii) 4

(iii) 774

(iv) 4196

#### **18.4.1a)**

**i)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)xl_1(B)w_1(B)u_1(B)u_1(A)xl_2(C)w_2(C)u_2(C)u_2(B)$   
1                      2                      3                      4    5

$xl_3(D)w_3(D)u_3(D)u_3(C)$   
6

**ii)**

$T_1$  is delayed. 4th lock request  $xl_1(B)$  waits since  $T_2$  is holding shared lock on B.

$T_2$  is delayed. 5th lock request  $xl_2(C)$  waits since  $T_3$  is holding shared lock on C.

$T_3$  gets  $xl_3(D)$ . Completes and unlocks C.

$T_2$  then can get  $xl_2(C)$ . Completes and unlocks B.

$T_1$  then can get  $xl_1(B)$  and completes.

**iii)**

Same as i)

**iv)**

Same as ii)

**v)**

Same as i)

**vi)**

Same as ii)

**18.4.1b)**

**i)**

$$\begin{array}{ccccccccc} sl_1(A) & r_1(A) & sl_2(B) & r_2(B) & sl_3(C) & r_3(C) & xl_1(B) & w_1(B) & u_1(B) & u_1(A) & xl_2(C) & w_2(C) & u_2(C) & u_2(B) \\ 1 & & 2 & & 3 & & 4 & & & & 5 & & & & \\ xl_3(A) & w_3(A) & u_3(A) & u_3(C) & & & & & & & & & & & \\ 6 & & & & & & & & & & & & & & \end{array}$$

**ii)**

T<sub>1</sub> delayed. 4<sup>th</sup> lock request xl<sub>1</sub>(B) waits since T<sub>2</sub> holds shared lock on B.  
T<sub>2</sub> delayed. 5<sup>th</sup> lock request xl<sub>2</sub>(C) waits since T<sub>3</sub> holds shared lock on C.  
T<sub>3</sub> delayed. 6<sup>th</sup> lock request xl<sub>3</sub>(A) waits since T<sub>1</sub> holds shared lock on A.  
Therefore, deadlock condition.

**iii)**

Same as i)

**iv)**

Same as ii)

**v)**

Same as i)

**vi)**

Same as ii)

**18.4.1c)**

**i)**

$$\begin{array}{ccccccccccc} sl_1(A) & r_1(A) & sl_2(B) & r_2(B) & sl_3(C) & r_3(C) & sl_1(B) & r_1(B) & sl_2(C) & r_2(C) & sl_3(D) & r_3(D) & xl_1(C) & w_1(C) & u_1(C) \\ 1 & & 2 & & 3 & & 4 & & 5 & & 6 & & 7 & & & \\ u_1(A) & u_1(B) & xl_2(D) & w_2(D) & u_2(D) & u_2(B) & u_2(C) & xl_3(E) & w_3(E) & u_3(E) & u_3(C) & u_3(D) & & & \\ & & 8 & & & & & & 9 & & & & & & \end{array}$$

**ii)**

T<sub>1</sub> is delayed. 7<sup>th</sup> lock request xl<sub>1</sub>(C) waits since T<sub>3</sub> and T<sub>2</sub> hold shared lock on C.

T<sub>2</sub> is delayed. 8<sup>th</sup> lock request  $xl_2(D)$  waits since T<sub>3</sub> holds shared lock on D.  
T<sub>3</sub> gets 9<sup>th</sup> lock request  $xl_3(E)$ . Completes. Unlocks D.  
T<sub>2</sub> then gets  $xl_2(D)$  and completes. Unlocks C and D.  
T<sub>1</sub> proceeds to completion.

**iii)**

Same as i)

**iv)**

Same as ii)

**v)**

Same as i)

**vi)**

Same as ii)

**18.4.1d)**

*i):  $xl_1(A); r_1(A); xl_2(B); r_2(B); xl_3(C); r_3(C); sl_1(B); r_1(B); sl_2(C); r_2(C); sl_3(D); r_3(D); w_1(A); u_1(A); u_1(B); w_2(B); u_2(B); u_2(C); w_3(C); u_3(C); u_3(D)$*

*(ii):* The first three locks and reads are allowed. However, T<sub>1</sub> cannot get a shared lock on B, nor can T<sub>2</sub> get a shared lock on C. When T<sub>3</sub> requests a shared lock on D it is granted. Thus, T<sub>3</sub> can proceed and eventually unlocks C and D.

As soon as C is unlocked, T<sub>2</sub> can get its shared lock on C. Thus, T<sub>2</sub> completes and releases its locks. As soon as its lock on B is released, T<sub>1</sub> can get its lock on A, so it completes.

*(iii):  $sl_1(A); r_1(A); sl_2(B); r_2(B); sl_3(C); r_3(C); sl_1(B); r_1(B); sl_2(C); r_2(C); sl_3(D); r_3(D); xl_1(A); w_1(A); u_1(A); u_1(B); xl_2(B); w_2(B); u_2(B); u_2(C); xl_3(C); w_3(C); u_3(C); u_3(D)$*

*(iv):* First, all six shared-lock requests are granted. When T<sub>1</sub> requests an exclusive lock on A, it gets it, because it is the only transaction holding a lock on A. T<sub>1</sub> completes and releases its locks. Thus, when T<sub>2</sub> asks for an exclusive lock on B, it is the only transaction still holding a shared lock on B, so that lock too may be granted, and T<sub>2</sub> completes. After T<sub>2</sub> releases its locks, T<sub>3</sub> is able to upgrade its shared lock on C to exclusive, and it too proceeds. The actions are thus executed in exactly the same order as they are requested; i.e., there are no delays by the scheduler.

*(v):  $ul_1(A); r_1(A); ul_2(B); r_2(B); ul_3(C); r_3(C); sl_1(B); r_1(B); sl_2(C); r_2(C); sl_3(D); r_3(D); xl_1(A); w_1(A); u_1(A); u_1(B); xl_2(B); w_2(B); u_2(B); u_2(C); xl_3(C); w_3(C); u_3(C); u_3(D)$*

(vi): The update locks prevent the fourth and fifth shared-lock requests,  $sl_1(B)$  and  $sl_2(C)$ , so  $T_1$  and  $T_2$  are delayed, while  $T_3$  is allowed to proceed. The situation is exactly as for part (ii), where the initial lock requests are for exclusive locks.

#### **18.4.1e)**

**i)**

$xl_1(A)r_1(A)xl_2(B)r_2(B)xl_3(C)r_3(C)sl_1(B)r_1(B)sl_2(C)r_2(C)sl_3(A)r_3(A)w_1(A)u_1(A)u_1(B)w_2(B)$   
 1                      2                      3                      4                      5                      6  
 $u_2(B)u_2(C)w_3(C)u_3(C)u_3(A)$

**ii)**

$T_1$  is delayed. 4<sup>th</sup> lock request  $sl_1(B)$  waits since  $T_2$  has exclusive lock on B.  
 $T_2$  is delayed. 5<sup>th</sup> lock request  $sl_2(C)$  waits since  $T_3$  has exclusive lock on C.  
 $T_3$  is delayed. 6<sup>th</sup> lock request  $sl_3(A)$  waits since  $T_1$  has exclusive lock on A.  
 Therefore, deadlock.

**iii)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)sl_1(B)r_1(B)sl_2(C)r_2(C)sl_3(A)r_3(A)xl_1(A)w_1(A)u_1(A)u_1(B)$   
 1                      2                      3                      4                      5                      6                      7  
 $xl_2(B)w_2(B)u_2(B)u_2(C)xl_3(C)w_3(C)u_3(C)u_3(A)$   
                     8    9

**iv)**

$T_1$  is delayed. 7<sup>th</sup> lock request  $xl_1(A)$  waits since  $T_3$  has shared lock on A.  
 $T_2$  is delayed. 8<sup>th</sup> lock request  $xl_2(B)$  waits since  $T_1$  has shared lock on B.  
 $T_3$  is delayed. 9<sup>th</sup> lock request  $xl_3(C)$  waits since  $T_2$  has shared lock on C.  
 Therefore, deadlock.

**v)**

$ul_1(A)r_1(A)ul_2(B)r_2(B)ul_3(C)r_3(C)sl_1(B)r_1(B)sl_2(C)r_2(C)sl_3(A)r_3(A)xl_1(A)w_1(A)u_1(A)$   
 1                      2                      3                      4                      5                      6                      7  
 $u_1(B)xl_2(B)w_2(B)u_2(B)u_2(C)xl_3(C)w_3(C)u_3(C)u_3(A)$   
                     8    9

**vi)**

$T_1$  is delayed. 4<sup>th</sup> lock request  $sl_1(B)$  waits since  $T_2$  has update lock on B.  
 $T_2$  is delayed. 5<sup>th</sup> lock request  $sl_2(C)$  waits since  $T_3$  has update lock on C.  
 $T_3$  is delayed. 6<sup>th</sup> lock request  $sl_3(A)$  waits since  $T_1$  has update lock on A.  
 Therefore, deadlock.

#### **18.4.2a)**

First, let us count the number of serializable orders that are conflict-equivalent to the order  $(T_1, T_2)$ . By symmetry, there will be the same number equivalent to the opposite

order,  $(T_2, T_1)$ . In order for a schedule to be equivalent to  $(T_1, T_2)$ ,  $inc_1(A)$  must precede  $r_2(A)$ . Thus, the first three steps of  $T_1$  must be first in the schedule.

Also,  $inc_1(B)$  must come before  $r_2(B)$ , so the former can either be the fourth action of the schedule, or it can be fifth, coming after  $r_2(A)$ . Thus, there are two serializable schedules equivalent to  $(T_1, T_2)$ , and four serializable schedules altogether.

#### **18.4.2b)**

Possible serializable schedules are

- (i) serial order  $(T_1, T_2)$
- (ii) serial order  $(T_2, T_1)$
- (iii) Serializable order equivalent to serial order  $(T_1, T_2)$  obtained by non-conflicting swaps.

Consider serial order  $(T_2, T_1)$ . For equivalent serializable schedule, the only non-conflicting swap allowed is between  $inc_1(B)$  and  $r_2(A)$ .

- (iv) Serializable order equivalent to serial order  $(T_2, T_1)$  obtained by non-conflicting swaps.

For serial order  $(T_2, T_1)$ , no non-conflicting swaps are possible since the last action of  $T_2$  i.e.  $inc_2(A)$  and first action of  $T_1$  i.e.  $r_1(A)$  are in conflict.

Thus total only three possible interleavings exist.

#### **18.4.3a)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)il_1(B)inc_1(B)il_2(C)inc_2(C)xl_1(C)w_1(C)u_1(A)u_1(B)u_1(C)xl_2(D)$   
 1                      2                      3                      4                      5    6  
 $w_2(D)u_2(B)u_2(C)u_2(D)$

$T_1$  is delayed. 3rd lock request  $il_1(B)$  waits since  $T_2$  has shared lock on B.

$T_2$  gets  $inc_2(C)$  and  $xl_2(D)$ .  $T_2$  completes and unlocks B, C and D.

$T_1$  then continues and gets  $il_1(B)$  and  $xl_1(C)$  and completes.

#### **18.4.3b)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)il_1(B)inc_1(B)il_2(A)inc_2(A)xl_1(C)w_1(C)u_1(C)u_1(A)u_1(B)xl_2(D)w_2(D)$   
 1                      2                      3                      4                      5    6  
 $u_2(D)u_2(B)u_2(A)$

$T_1$  is delayed. 3<sup>rd</sup> lock request  $il_1(B)$  waits since  $T_2$  has shared lock on B.

$T_2$  is delayed. 4<sup>th</sup> lock request  $il_2(A)$  waits since  $T_1$  has shared lock on A.

Therefore, deadlock.

#### **18.4.3c)**

$il_1(A)inc_1(A)il_2(B)inc_2(B)il_1(B)inc_1(B)il_2(C)inc_2(C)xl_1(C)w_1(C)u_1(C)u_1(B)u_1(A)xl_2(D)$



1                      2                      3                      4                      5                      6  
 $w_2(D)u_2(D)u_2(B)u_2(C)$

$il_1(B)$  does not conflict with  $il_2(B)$ .

$T_1$  is delayed. 5<sup>th</sup> lock request  $xl_1(C)$  waits since  $T_2$  holds increment lock on C.

$T_2$  completes and unlocks C.

$T_1$  then gets  $xl_1(C)$  and completes.

#### **18.4.4**

Schedule is  $r_1(A)r_1(B)w_1(B)r_1(C)w_1(C)r_1(D)w_1(D)r_1(E)w_1(E)$

With locks:

$Ul_1(A)r_1(A)ul_1(B)r_1(B)xl_1(B)w_1(B)ul_1(C)r_1(C)xl_1(C)w_1(C)ul_1(D)r_1(D)xl_1(D)w_1(D)ul_1(E)r_1(E)xl_1(E)w_1(E)$

(Quest: shared or update lock better for concurrency?)

#### **18.4.5a)**

		Requested		
Held		<b>S</b>	<b>X</b>	<b>M</b>
	<b>S</b>	Yes	No	No
	<b>X</b>	No	No	No
	<b>M</b>	No	No	Yes

#### **18.4.5b)**

		Requested			
Held		<b>S</b>	<b>X</b>	<b>I</b>	<b>M</b>
	<b>S</b>	Yes	No	No	No
	<b>X</b>	No	No	No	No
	<b>I</b>	No	No	Yes	No
	<b>M</b>	No	No	No	Yes

#### **18.4.6a)**

The only actions that commute are (i)-(ii) and (iii)-(iv). That is, it doesn't matter whether we first set the x-axis and then the y-axis, or vice-versa. Likewise, it doesn't matter whether we first set the angle and then the magnitude or vice-versa. However, all other pairs of actions matter.

For example, consider the actions (set angle = 90)(set x = 5) applied to the vector (10,0). The first step gives us (0,10), and the second step gives us (5,10).

On the other hand, the reverse order of the actions: (set x = 5)(set angle = 90) starting with (10,0) gives us (0,5), and the second step gives us (5,5).

**18.4.6b)**

		Requested			
Held		X	Y	A	M
	X	No	Yes	No	No
	Y	Yes	No	No	No
	A	No	No	No	Yes
	M	No	No	Yes	No

**18.4.6c)**

		Requested			
Held		X	Y	A	M
	X	Yes	Yes	No	No
	Y	Yes	Yes	No	No
	A	No	No	Yes	Yes
	M	No	No	Yes	Yes

**18.4.7a)**

The only constraint on a serial order so far is that  $T_1$  must precede  $T_2$ , because  $T_1$  reads  $A$  before  $T_2$  writes  $A$ . The only way we could get a cycle in the precedence graph is if ??? required  $T_2$  to precede  $T_1$ . A read of anything by  $T_1$  will not introduce any constraints, but  $r_2(C)$  will cause an arc from  $T_2$  to  $T_1$  and lead to a cycle.

**18.4.7b)**

$w_1(B)$  will require  $T_2$  to precede  $T_1$ . Given that  $T_1$  is required to precede  $T_1$  since  $T_1$  reads  $A$  before  $T_2$  writes  $A$ , a cycle is introduced in precedence graph.

**18.4.7c)**

$r_1(C)$  or  $r_2(A)$  will cause update locks.

Both of these actions do not lead to a cycle. So, no schedule is non-serializable.

**18.4.7d)**

$inc_2(C)$  will conflict with  $w_1(C)$ . It will require  $T_2$  to precede  $T_1$  leading to a cycle.

Also,  $incr_1(B)$  will conflict with  $r_2(B)$  requiring  $T_2$  to precede  $T_1$ , thus causing a cycle.

**18.5.1a)**

Based on the compatibility matrix of shared and exclusive locks, the only possible sets of locks on a single element are either:

- i) One exclusive lock
- ii) One or more shared locks

Thus, the only needed lock modes are exclusive and shared, which correctly sum up the limitations on access to a given element in the above two situations, respectively.

### **18.5.1b)**

Based on the compatibility matrix of shared, exclusive and increment locks, we see that the only possible sets of locks on a single element are either

- i) One exclusive lock.
- ii) One or more shared locks.
- iii) One or more increment locks.

Thus, the only needed lock modes are exclusive, shared, and increment, which correctly sum up the limitations on access to a given element in the above three situations, respectively.

### **18.5.1c)**

Based on the compatibility matrix drawn up in 18.4.6b, the only possible sets of locks on a single element are either:

- i) One X-Lock and one Y-lock
- ii) One A-lock and one M-lock
- iii) One X-lock
- iv) One Y-lock
- v) One M lock
- vi) One A lock

Each of the above six situations represent a lock mode.

### **18.5.2a)**

$r_1(A)r_2(A)r_3(B)w_1(A)r_2(C)r_2(B)w_2(B)w_1(C)$

The lock scheduler described in this section inserts shared-exclusive-update locks. Thus the sequence of lock requests and actions for the above schedule is:

$u_1(A)r_1(A)sl_2(A)r_2(A)sl_3(B)r_3(B)xl_1(A)w_1(A)sl_2(C)r_2(C)ul_2(B)r_2(B)xl_2(B)w_2(B)xl_1(C)w_1(C)$   
)

This is how the lock scheduler would execute:

- $T_1$  issues  $r_1(A)$ . Part 1 of scheduler inserts  $u_1(A)$  since there is a future upgrade to exclusive lock needed. Part 2 of scheduler consults the lock table and grants  $u_1(A)$  since there is no other lock on A.

- Next  $T_2$  issues  $r_2(A)$ . Part 1 inserts  $sl_2(A)$  since there is no future upgrading to exclusive lock. Part 2 consults lock table and finds that it cannot grant  $sl_2(A)$  since  $T_1$  holds an update lock to A. So, this and the subsequent actions from  $T_2$  are delayed.
- $T_3$  then issues  $r_3(B)$ . Part 1 insert  $sl_3(B)$ . Part 2 grants the shared lock on B to  $T_3$ .  $T_3$  completes action and commits. Part1 then releases lock on B.
- $T_1$  issues  $w_1(A)$ . Part1 inserts an exclusive lock  $xl_1(A)$ . Part 2 upgrades the lock to exclusive.
- $T_1$  issues  $w_1(C)$ . Part 1 issues  $xl_1(C)$ . Part 2 grants the lock.  $T_1$  commits and Part1 releases locks on A and C.
- Its then found out that T is waiting for lock on A. Part 2 is notified and it grants  $sl_2(A)$ .  $T_2$  then completes.

### **18.5.2b)**

$r_1(A)w_1(B)r_2(B)w_2(C)r_3(C)w_3(A)$

The sequence of lock requests in the schedule (assuming no delays) would be:

$sl_1(A)r_1(A)xl_1(B)w_1(B)sl_2(B)r_2(B)xl_2(C)w_2(C)sl_3(C)r_3(C)xl_3(A)w_3(A)$

The lock scheduler would execute as follows:

- $T_1$  issues  $r_1(A)$ . Part 1 inserts  $sl_1(A)$  since no future upgrades are required. Part 2 grants  $sl_1(A)$  after consulting the lock table.
- $T_1$  issues  $w_1(B)$ . Part1 inserts  $xl_1(B)$ . Part 2 grants  $xl_1(B)$  and updates the lock table.
- $T_1$  commits and Part 1 then unlocks A and B.
- $T_2$  issues  $r_2(B)$ . Part 1 inserts  $sl_2(B)$ . Part 2 grants.
- $T_2$  issues  $w_2(C)$ . Part1 inserts  $xl_2(C)$ . Part 2 grants.
- $T_2$  commits. Part 1 releases locks on B and C.
- Similarly,  $T_3$  completes. So this schedule completes without any delays.

### **18.5.2c)**

$w_3(A)r_1(A)w_1(B)r_2(B)w_2(C)r_3(C)$

Here we show the sequence of lock requests inserted by Part 1 and granted by Part 2 after consulting the lock table. We do not provide detailed explanation since there are no delays in this schedule.

$xl_3(A)w_3(A)sl_1(A)r_1(A)xl_1(B)w_1(B)T_1$  commits.  $ul_1(A)ul_1(B)sl_2(B)r_2(B)sl_2(C)w_2(C)$ .  $T_2$  commits.  $ul_2(C)$ .  $ul_2(B)sl_3(C)r_3(C)$ .  $T_3$  commits.  $Ul_3(C)$ .

### **18.5.2d)**

$r_1(A)r_2(A)w_1(B)w_2(B)r_1(B)r_2(B)w_2(C)w_1(D)$

Here we show the sequence of lock requests (assuming no delays):

$sl_1(A)r_1(A)sl_2(A)r_2(A)xl_1(B)w_1(B)xl_2(B)w_2(B)r_1(B)sl_2(B)r_2(B)xl_2(C)w_2(C)ul_2(A)ul_2(B)ul_2(C)xl_1(D)w_1(D)ul_1(A)ul_1(B)ul_1(D)$

Here is how the lock scheduler would execute:

- $T_1$  issues  $r_1(A)$ . Part 1 issues and Part 2 grants  $sl_1(A)$ . Lock table is updated.

- $T_2$  issues  $r_2(A)$ . Part 1 inserts  $sl_2(A)$ . Part 2 consults the lock table and it sees that the only lock held on A is a shared lock and thus it grants  $sl_2(A)$ .
- $T_1$  issues  $w_1(B)$ . Part1 inserts  $xl_1(B)$  and Part 2 issues it since there are no locks on B.
- $T_2$  issues  $w_2(B)$ . Part 2 issues  $xl_2(B)$ . Part 2 consults the lock table for B and finds out that  $T_1$  already holds an exclusive lock on B.  $T_2$  is thus delayed.
- $T_1$  then issues  $r_1(B)$ . Since  $T_1$  already holds an exclusive lock on B, no new lock request is inserted and  $T_1$  performs  $r_1(B)$ .
- $T_1$  issues  $w_1(D)$ . Part 1 inserts and Part 2 grants  $xl_1(D)$ .
- $T_1$  then commits and Part 1 releases locks on A, B and D.
- Part 2 is then notified that  $T_2$  is waiting on  $xl_2(B)$ . It grants  $xl_2(B)$ .
- $T_2$  then completes.

### **18.5.2e)**

$r_1(A)r_2(A)r_1(B)r_2(B)r_3(A)r_4(B)w_1(A)w_2(B)$

Here we show the sequence of lock requests (assuming no delays):


$ul_1(A)r_1(A)sl_2(A)r_2(A)sl_1(B)r_1(B)ul_2(B)r_2(B)sl_3(A)r_3(A)sl_4(B)r_4(B)xl_1(A)w_1(A)ul_1(A)ul_1(B)$

Here is how the lock scheduler would execute:

- $T_1$  issues  $r_1(A)$ . Part 1 inserts  $ul_1(A)$ . Part 2 grants  $ul_1(A)$ .
- $T_2$  issues  $r_2(S)$ . Part 1 inserts  $sl_2(A)$ . Part 2 consults the lock table and denies  $sl_2(A)$  since lock table shows  $T_1$  holds an update lock to A. The actions of  $T_2$  are thus delayed.
- $T_1$  issues  $r_1(B)$ . Part 2 issues  $sl_1(B)$ .
- $T_3$  issues  $r_3(A)$ . Part 1 inserts  $sl_3(A)$ . Part 2 consults the lock table and denies  $sl_3(A)$  since  $T_1$  still holds an update lock to A.  $T_3$  is thus delayed.
- $T_4$  issues  $r_4(B)$ . Part 1 inserts  $sl_4(B)$ . Part 2 consults the lock table for B and denies  $sl_4(B)$  since  $T_1$  holds  $sl_1(B)$ .  $T_4$  is thus delayed.
- $T_1$  issues  $w_1(A)$ . Part 2 upgrades  $T_1$ 's lock on A to  $xl_1(A)$ .
- $T_1$  commits and Part 1 releases locks on A and B.
- Part 2 is notified that  $T_2$  is waiting on  $Sl_2(A)$ . Part 2 grants  $sl_2(A)$ .
- $T_2$  issues  $r_2(B)$ . Part 1 issues and Part 2 grants  $ul_2(B)$ .
- $T_3$  and  $T_4$  still wait.
- $T_2$  issues  $w_2(B)$ . Part 2 upgrades  $T_2$ 's lock to  $xl_2(B)$ .
- $T_2$  then completes and Part 1 releases locks on A and B.
- $T_3$  and  $T_4$ , which are waiting on A and B respectively, get the locks and complete.

### **18.6.1a)**

Refer to the following hierarchy for all four sub-parts:



- i) T<sub>1</sub> puts IS lock on C and B<sub>1</sub>, and S lock on O<sub>1</sub>
- ii) T<sub>2</sub> puts IX lock on C, IX lock on B<sub>1</sub> and X lock on O<sub>2</sub>
- iii) T<sub>2</sub> puts IS lock on C, B<sub>2</sub> and S lock on O<sub>3</sub>. T<sub>2</sub> releases its locks.
- iv) T<sub>1</sub> puts IX lock on C, B<sub>2</sub> and S lock on O<sub>4</sub>. T<sub>1</sub> releases its locks.

#### **18.6.1b)**

- i) T<sub>1</sub> puts IS lock on C, B<sub>2</sub> and S lock on O<sub>5</sub>.
- ii) T<sub>2</sub> puts IX lock on C, B<sub>2</sub>. It attempts X lock on O<sub>5</sub>, which is rejected since T<sub>1</sub> already holds an S lock on O<sub>5</sub>. T<sub>2</sub> is delayed.
- iii) T<sub>1</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>4</sub>. T<sub>1</sub> releases all locks
- iv) T<sub>2</sub> then puts X lock on O<sub>5</sub>.
- v) T<sub>2</sub> puts IS lock on C, B<sub>2</sub> and S lock on O<sub>3</sub>. T<sub>2</sub> releases all locks.

#### **18.6.1c)**

- i) T<sub>1</sub> puts IS lock on C, B<sub>1</sub> and S lock on O<sub>1</sub>
- ii) T<sub>1</sub> puts IS lock on C, B<sub>2</sub> and S lock on O<sub>3</sub>. T<sub>1</sub> releases locks.
- iii) T<sub>2</sub> puts IS lock on C, B<sub>1</sub> and S lock on O<sub>1</sub>
- iv) T<sub>2</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>4</sub>
- v) T<sub>2</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>5</sub>. T<sub>2</sub> releases locks.

#### **18.6.2d)**


- i) T<sub>1</sub> puts IS lock on C, B<sub>1</sub> and S lock on O<sub>1</sub>
- ii) T<sub>2</sub> puts IS lock on C, B<sub>1</sub> and S lock on O<sub>2</sub>
- iii) T<sub>3</sub> puts IS lock on C, B<sub>1</sub> and S lock on O<sub>1</sub>
- iv) T<sub>1</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>3</sub>
- v) T<sub>2</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>4</sub>. T<sub>2</sub> releases locks.
- vi) T<sub>3</sub> puts IX lock on C, B<sub>2</sub> and X lock on O<sub>5</sub>. T<sub>3</sub> releases locks.
- vii) T<sub>1</sub> puts IX lock on C, B<sub>1</sub> and X lock on O<sub>2</sub>. T<sub>1</sub> releases locks.

#### **18.6.2**

By replacing w<sub>4</sub>(D<sub>3</sub>) by w<sub>4</sub>(Movie), the schedule becomes:

R<sub>3</sub>(D<sub>1</sub>)r<sub>3</sub>(D<sub>2</sub>)w<sub>4</sub>(Movie)w<sub>4</sub>(X)w<sub>3</sub>(L)w<sub>3</sub>(X)

The hierarchy of elements:



The sequence of actions of warning protocol scheduler are:

- i) T3 puts IS lock on Movie and S lock on D1.
- ii) T3 puts IS lock on Movie and S lock on D2.
- iii) T4 attempts X lock on Movie, but is rejected since X-lock and IS lock are not compatible. T4 is delayed.
- iv) T3 computes  $L = \text{sum of lengths of D1 and D2}$ .
- v) T3 puts IX lock on Movie and X lock on X. T3 releases locks.
- vi) T4 then gets X lock on Movie. Updates movie. Then gets X lock on X and updates X. T4 releases locks.

The sequence of actions is equivalent to serial order (T3,T4).

### **18.6.3**

With increment locks, the steps of the warning protocol scheduler stay the same, except the compatibility matrix is updated to include II (intention increment) and I (increment) locks as following:

	IS	IX	II	S	X	I
IS	Y	Y	Y	Y	N	N
IX	Y	Y	Y	N	N	N
II	Y	Y	Y	N	N	Y
S	Y	N	N	Y	N	N
X	N	N	N	N	N	N
I	N	N	Y	N	N	Y

### **18.7.1a)**

As soon as we reach the left child of the root, we see that node is not full. Thus, the insert of 10 cannot cause that node to split, and there will be no reason to rewrite the root. We can now release the exclusive lock on the root.

We are then directed to the second leaf. Since that leaf is not full either, we know the insertion will not cause the left child of the root to be changed, so we can release the exclusive lock on that child. As soon as we have rewritten the second leaf to reflect the insertion of 10, we can release the lock on the leaf.

### **18.7.1b)**

Lookup of 20 leads to right child of root node. Since right child is full, we might need to rewrite root, so do not release exclusive lock on root. Now,  $20 < 23$ , so get directed to third leaf, which is full. Therefore, it needs to be split. We also do not release the exclusive lock on right child of root since the split of leaf will cause the right child to be updated. First the exclusive lock on third leaf gets released after rewrite, followed by right child and followed by root node.

#### **18.7.1c)**

For delete of 5, lookup of 5 leads to left child of root node. We then get directed to first leaf since  $5 < 7$ . We see that delete of 5 would not cause unbalanced leaf. So we can release the exclusive lock on left child and root node before doing the delete operation.

#### **18.7.1d)**

Lookup of 23 leads us to the right child of root node. We can release the exclusive lock on root node since delete of 23 would not cause root to be updated. We follow the second pointer in the right child and reach the fourth leaf. Delete of 23 would cause the fourth leaf to have only 1 key, which would cause 19 to be moved from third to fourth leaf. We thus hold on to the exclusive lock on second child of root until after updating the leafs since the second child needs to be updated as well.

#### **18.7.2a)**

Note: we shall assume only a single type of lock, so that even though the operations are all reads (and therefore, any interleaving could in principle be considered serializable), it is not possible for two transactions to access the same element until one has relinquished its lock.

Suppose that  $T_1$  locks  $A$  first. Then  $T_2$  cannot perform any steps until  $T_1$  relinquishes the lock. Therefore,  $T_1$  must also read  $B$ , at the second step of the schedule. Now,  $T_1$  can relinquish the lock on  $A$ . The remaining step of  $T_1$ ,  $r_1(E)$ , can occur either before all of  $T_2$ , or after the first or second step of  $T_2$ . It cannot occur after the last step of  $T_2$ ,  $r_2(B)$ , because  $T_1$  can't relinquish its lock on  $B$ , until it has locked  $E$ . Thus, there are three interleavings in which  $T_1$  starts first.

Now, consider what happens if  $T_2$  starts first.  $T_2$  cannot relinquish its lock on  $A$  until it has locked both children  $B$  and  $C$ . Thus, if  $T_2$  starts first, it must finish before  $T_1$  starts. We conclude that there are four legal schedules.

#### **18.7.2b)**

$T_1$  can lock  $A$  anytime. Next  $T_1$  can get a lock on  $B$  if  $T_3$  didn't lock  $B$  yet. If  $T_1$  gets a lock on  $B$  first it will get a lock on  $E$  first also. Now  $T_1$  can release a lock on  $B$  and  $T_3$  can get it but  $T_3$  also need to wait to get a lock on  $E$  until  $T_1$  releases it. After getting a lock on  $E$ ,  $T_3$  can lock  $F$ . Thus, if  $T_1$  locks  $B$  first,  $T_3$  can start after  $T_1$  gets a lock and  $E$  and release a lock on  $B$  and  $T_3$  can finish all steps after  $T_1$  finishes all steps. The same situation occurs when  $T_3$  locks  $B$  first. Therefore, there are two interleavings.

1)  $T_1 \rightarrow T_3$ :



$$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E) \rightarrow$$

$$r_3(B) \rightarrow r_3(E) \rightarrow r_3(F),$$

If T1 reads B first, T2 needs to wait to get a lock on B until T1 gets a lock on E and releases a lock on B.

2) T3  $\rightarrow$  T1:

$$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$$

$$r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$$

If T2 locks B first, T1 needs to get a lock on B until T2 gets a lock on E and releases a lock on B.  $r_1(A)$  can start in anytime either before all steps, between steps, or after all steps (4 cases) and  $r_3(F)$  can start either before or after  $r_1(B)$  but not after  $r_1(E)$  (2 cases). Therefore, all (3 cases  $r_1(A)$  \* 2 cases of  $r_1(E)$ ) + 1 case of  $r_1(A)$  starts after  $r_3(F)$  = 7 interleavings can occur.

Thus, all 8 possible interleavings here.

### **18.7.2c)**

1) T1 – T2 – T3 : ... The sequence of Ts get locks on a common element.

$$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$$

$$r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$$

$$r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$$

If T3 gets a lock on B first, T3 can start either before or after  $r_2(C)$  (2 cases) and  $r_2(B)$  can start before or after  $r_3(F)$  (2 cases) ( $2 * 2 = 4$  cases). Thus, there are all 5 interleavings here.

2) T1 – T3 – T2 :

$$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$$

$$r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$$

$$r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$$

If T2 gets a lock on B first, T3 needs wait until T2 finishes all steps. Thus, there is only one interleaving here.

3) T2 – T1 – T3 :

$$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$$

$r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$   
 $r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$

There's only one interleaving.

4) T2 – T3 – T1 :

$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$   
 $r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$   
 $r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$

$r_1(A)$  can start before  $r_2(B)$ , after  $r_3(F)$ , or between those two (5 cases).  $r_1(B)$  can start before or after  $r_3(F)$  (2 cases). (4 cases of  $r_1(A)$  \* 2 cases of  $r_1(B)$ ) + 1 cases that  $r_1(A)$  starts after  $r_3(F)$ . Thus, 9 interleavings here.

5) T3 – T1 – T2 :

$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$   
 $r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$   
 $r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$

$r_1(A)$  can start anytime before, between, or after all steps of T3 (4 cases) and  $r_3(F)$  can start before or after  $r_1(B)$  but not after  $r_1(E)$  (2 cases).  $r_1(E)$  can start anytime in T2 but not after  $r_2(B)$  (3 cases). Thus, there are 24 interleavings from this sequence.

6) T3 – T2 – T1:

$r_1(A) \rightarrow r_1(B) \rightarrow r_1(E)$   
 $r_2(A) \rightarrow r_2(C) \rightarrow r_2(B)$   
 $r_3(B) \rightarrow r_3(E) \rightarrow r_3(F)$

$r_2(A)$  can start before, after, or between all steps of T3 (4 cases). So can  $r_2(C)$ . But  $r_2(B)$  can start only before or after  $r_3(F)$ .  $r_1(A)$  can only start before or after  $r_2(B)$ .  $(4 * 4 - 1(\text{the case both } r_2(A) \text{ and } r_2(C) \text{ starts after } r_3(F)) * 2) = 30$  interleavings.

Therefore, all possible interleavings are  
 $5 + 1 + 1 + 9 + 16 + 24 + 30 = 86$

### **18.7.3**

$R(T1-T3-T5-T7)$   
 $|$   
 $+-----+-----+$   
 $A(T1-T2-T3-T4) \quad B(T3-T6-T5) \quad C(T8-T7)$   
  
 $T1 \quad T2 \quad T3 \quad T4 \quad T5 \quad T6 \quad T7 \quad T8$   
 $-----$

```

l1(R); r1(R);
l8(C); r8(C);
l1(A); r1(A);
u8(C);
u1(R);          l3(R); r3(R);
u1(A)          l2(A); r2(A); l3(B); r3(B);
                u2(A);          l3(A); r3(A);
                                u3(B); u3(R);
                                u3(A);          l4(A); r4(A);
                                                l5(R); r5(R); l6(B); r6(B);
                                                u5(B); u5(R);          u6(B);
                                                                l7(R); r7(R);
                                                                l7(C); r7(C);
                                                                u7(R); u7(C);

```

---

$T1 < T2 < T3 < T4$   
 $T3 < T6 < T5 < T7$   
 $T8 < T7$


T4 and T8 are not dependent on any other transaction after T3 and before T7 respectively.  
 4 possible serial order based on T4's position (before T6, after T7, or between those) \* 6  
 possible serial order based on T8's position (before T1, after T5, or between those) = 24  
 serial orders are consistent with the statement.

#### 18.7.4

Let's consider the following serializable transaction.

S: w1(C); r1(B); r2(B); w2(D)

Suppose we have the following tree.



Based on Rule(2) of the tree protocol, each transaction needs to get a lock on a parent node to get a shared or exclusive lock of the subsequent node. It can lead S', which is not serializable, when T2 gets a exclusive lock to get a exclusive lock on D before T1 whether T1 starts r2(B) first or w2(B) first.

S': w1(A); w1(C); w2(B); r2(B); w2(D); r1(B)

Change Rule (2) to "Subsequent shared or exclusive lock may be acquired if the transaction has a **shared lock** on its parent node" can prevent unserializable behavior and it leads S'' based on the changed Rule(2).

S'': r1(A); w1(C); r2(B); w2(D); r1(B),

which can be converted into

S'':  $r_1(A)$ ;  $w_1(C)$ ;  $r_1(B)$ ;  $r_2(B)$ ;  $w_2(D)$ ;

**18.8.1a)**

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>A</u>	<u>B</u>
		RT = 0 WT=0	RT=0 WT=0
st <sub>1</sub> Assume TS(T <sub>1</sub> )=1			
	st <sub>2</sub> Assume TS(T <sub>2</sub> )=2		
r <sub>1</sub> (A)		RT=1	
	r <sub>2</sub> (B)		RT=2
	w <sub>2</sub> (A)	WT=2	
w <sub>1</sub> (B) <b>T<sub>1</sub> aborts since TS(t<sub>1</sub>)&lt;RT(B)</b>			

**18.8.1b)**

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>A</u>	<u>B</u>
		RT=0 WT=0	RT=0 WT=0
st <sub>1</sub> Assume TS(T <sub>1</sub> )=1			
r <sub>1</sub> (A)		RT=1	
	st <sub>2</sub> Assume TS(T <sub>2</sub> )=2		
	w <sub>2</sub> (B)		WT=2
	r <sub>2</sub> (A)	RT=2	
w <sub>1</sub> (B) TS(T <sub>1</sub> )<WT(B) i.e. later trans has already written B. So w <sub>1</sub> (B) ignored.			

**18.8.1c)**

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>T<sub>3</sub></u>	<u>A</u>	<u>B</u>	<u>C</u>
----------------------	----------------------	----------------------	----------	----------	----------

			RT=0 WT=0	RT=0 WT=0	RT=0 WT=0
st <sub>1</sub> . Assume TS(T <sub>1</sub> ) = 1					
	st <sub>2</sub> . Assume TS(T <sub>2</sub> )=2				
		st <sub>3</sub> . Assume TS(T <sub>3</sub> )=3			
r <sub>1</sub> (A)			RT=1		
	r <sub>2</sub> (B)			RT=2	
w <sub>1</sub> (C)					WT=1
		r <sub>3</sub> (B)		RT=3	
		r <sub>3</sub> (C)			RT=3
	w <sub>2</sub> (B) TS(T <sub>2</sub> )<RT(B) <b>T2 aborts.</b>				
		w <sub>3</sub> (A)	WT=3		

#### 18.8.1d)

<u>T1</u>	<u>T2</u>	<u>T3</u>	<u>A</u>	<u>B</u>	<u>C</u>
			RT=0 WT=0	RT=0 WT=0	RT=0 WT=0
st <sub>1</sub> . Assume TS(T <sub>1</sub> )=1					
		st <sub>3</sub> . Assume TS(T <sub>3</sub> )=2			
	st <sub>2</sub> . Assume TS(T <sub>2</sub> )=3				
r <sub>1</sub> (A)			RT=1		
	r <sub>2</sub> (B)			RT=3	
w <sub>1</sub> (C)					WT=1
		r <sub>3</sub> (B) No change in RT(B) since RT(B)>TS(T <sub>3</sub> )		RT=3	
		r <sub>3</sub> (C)			RT=2
	w <sub>2</sub> (B) Since TS(T <sub>2</sub> )=RT(B), unchanged			WT=3	
		W <sub>3</sub> (A)	WT=2		

#### 18.8.2a)

The three writes create three versions of *A*. When T<sub>2</sub> tries to read *A*, it is given the value that it itself wrote, since that is the version with the greatest timestamp that does not

exceed the timestamp of  $T_2$ . That makes sense, although in practice, we doubt that a well written transaction would read its own value through the database storage system.

When  $T_4$  tries to read  $A$  the system finds that  $T_4$ 's timestamp is larger than that of any version of  $A$  written. Thus,  $T_4$  gets the version with the largest of the timestamps, the one written by  $T_3$ . That makes sense, because in the hypothetical serial order based on the timestamps of the transactions,  $T_3$  would be the last to write  $A$ .

#### 18.8.2b)


<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>T<sub>3</sub></u>	<u>T<sub>4</sub></u>	<u>A<sub>1</sub></u>	<u>A<sub>3</sub></u>
st1. Assume TS(T <sub>1</sub> )=1	st2. Assume TS(T <sub>2</sub> )=2	st3. Assume TS(T <sub>3</sub> )=3	st4. Assume TS(T <sub>4</sub> )=4		
w <sub>1</sub> (A)				Create	
		w <sub>3</sub> (A)			Create
			r <sub>4</sub> (A)		Read
	r <sub>2</sub> (A)			Read	

Without multi-version timestamp scheduler,  $T_2$  would have aborted on issuing  $r_2(A)$  since  $TS(T_2)$  [2] would be less than write time of  $A$  [3].

#### 18.8.2c)

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>	<u>T<sub>3</sub></u>	<u>T<sub>4</sub></u>	<u>A<sub>1</sub></u>	<u>A<sub>4</sub></u>
st1. Assume TS(T <sub>1</sub> )=1	st2. Assume TS(T <sub>2</sub> )=2	st3. Assume TS(T <sub>3</sub> )=3	st4. Assume TS(T <sub>4</sub> )=4		
w <sub>1</sub> (A)				Create	
			w <sub>4</sub> (A)		Create
		r <sub>3</sub> (A)		Read	
	w <sub>2</sub> (A) <b>T<sub>2</sub> aborts.</b>				

On  $w_2(A)$ , following occurs:



Without multi-version timestamp scheduler,  $T_3$  would have aborted on  $r_3(A)$  since  $A$  would have already been written by a later transaction  $T_4$ . Also,  $w_2(A)$  would have been ignored since  $T_4$  with a later timestamp would have written  $A$ .

### **18.8.3**

When a younger transaction  $T_2$  requests an item  $X$  held by an older transaction  $T_1$ ,  $T_2$  waits until the commit bit,  $C(X)$ , becomes true.  $T_1$  never waits for  $T_2$ . Wait-die strategy is for  $T_2$ , where the scheduler let  $T_2$  be aborted and started again when  $T_2$  waits too long.

When an older transaction  $T_1$  requests an item  $X$  held by a younger transaction  $T_2$ ,  $T_1$  waits for  $T_2$ .  $T_2$  never waits for  $T_1$ . Wound-wait strategy is for  $T_1$ , where the scheduler preempts  $T_2$  by aborting it.

In either situation, no cycle is introduced between  $T_1$  and  $T_2$ . Thus, no deadlock situation occurs.

### **18.9.1a)**

As  $T_1$  is the first to validate, there is nothing to check;  $T_1$  validates successfully.

$T_3$  validates next. The only other validated transaction is  $T_1$ , and  $T_1$  has not yet finished. Thus, both the read- and write-sets of  $T_3$  must be compared with the write-set of  $T_1$ . However,  $T_1$  writes only  $A$ , and  $T_3$  neither reads nor writes  $A$ , so  $T_3$ 's validation succeeds.

Last,  $T_2$  validates. Both  $T_1$  and  $T_3$  finish after  $T_2$  started, so we must compare the read-set of  $T_2$  with the write-sets of both  $T_1$  and  $T_3$ . Since  $B$  is in both  $W_3$  and  $R_2$ , we cannot validate  $T_2$ . Note that since  $T_3$  (but not  $T_1$ ) finishes after  $T_2$  validates, we would also compare the write set of  $T_2$  with the write set of  $T_3$ , had we not already found a reason not to validate  $T_2$ .

### **18.9.1b)**

The sequence of validation steps are explained below:

- $v_1$ 
  - No transaction has validated yet. So,  $T_1$  validates.
- $v_3$  ( $T_1$  has already validated)
  - $FIN(T_1) > START(T_3)$  since when  $T_3$  started,  $T_1$  has not finished.  $WS(T_1) \cap RS(T_3). (A) \cap (C,D) = \emptyset$ .
  - $FIN(T_1) > VAL(T_3)$ .  $WS(T_1) \cap WS(T_3). (A) \cap (D) = \emptyset$ .  $T_3$  validates.
- $v_2$  ( $T_1$  and  $T_3$  have validated)
  - $FIN(T_1) > START(T_2)$ .  $RS(T_2) \cap WS(T_1). (B,C) \cap A = \emptyset$ .
  - $FIN(T_3) > START(T_2)$ .  $RS(T_2) \cap WS(T_3). (B,C) \cap D = \emptyset$ .
  - $FIN(T_3) > VAL(T_2)$ .  $WS(T_2) \cap WS(T_3). A \cap D = \emptyset$ .  $T_2$  validates.

### **18.9.1c)**

- $v_1$ 
  - No transaction validated yet. So  $T_1$  validates.

- $v_3$  ( $T_1$  has validated)
  - $\text{FIN}(T_1) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_1). (C,D) \cap (C) = C$ .  $T_3$  rolled back.
- $v_2$  ( $T_1$  has validated)
  - $\text{FIN}(T_1) > \text{START}(T_2). \text{RS}(T_2) \cap \text{WS}(T_1). (B,C) \cap (C) = C$ .  $T_2$  also rolled back.

#### **18.9.1.d)**

- $v_1$ 
  - No transaction validated yet. So  $T_1$  validates.
- $v_2$  ( $T_1$  already validated)
  - $\text{FIN}(T_1) > \text{START}(T_2). \text{RS}(T_2) \cap \text{WS}(T_1). (B,C) \cap A = \emptyset$ .
  - $\text{FIN}(T_1) > \text{VAL}(T_2). \text{WS}(T_1) \cap \text{WS}(T_2). A \cap B = \emptyset$ .  $T_2$  validates.
- $v_3$  ( $T_1$  and  $T_2$  have validated)
  - $\text{FIN}(T_1) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_1). C \cap A = \emptyset$ .
  - $\text{FIN}(T_1) > \text{VAL}(T_3). \text{WS}(T_3) \cap \text{WS}(T_1). A \cap C = \emptyset$ .
  - $\text{FIN}(T_2) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_2). C \cap B = \emptyset$ .
  - $\text{FIN}(T_2) > \text{VAL}(T_3). \text{WS}(T_3) \cap \text{WS}(T_2). B \cap C = \emptyset$ .  $T_3$  validates.

#### **18.9.1e)**

- $v_1$ 
  - No transaction validated yet.  $T_1$  validates.
- $v_2$  ( $T_1$  validated)
  - $\text{FIN}(T_1) > \text{START}(T_2). \text{RS}(T_2) \cap \text{WS}(T_1). (B,C) \cap (C) = C$ .  $T_2$  rolled back.
- $v_3$  ( $T_1$  validated)
  - $\text{FIN}(T_1) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_1). (C) \cap (C) = C$ .  $T_3$  rolled back.

#### **18.9.1f)**

- $v_1$ 
  - No transaction validated yet.  $T_1$  validates.
- $v_2$  ( $T_1$  is validated)
  - $\text{FIN}(T_1) > \text{START}(T_2). \text{RS}(T_2) \cap \text{WS}(T_1). (B,C) \cap (A) = \emptyset$ .
  - $\text{FIN}(T_1) > \text{VAL}(T_2). \text{WS}(T_2) \cap \text{WS}(T_1). A \cap C = \emptyset$ .  $T_2$  validates.
- $v_3$  ( $T_1$  and  $T_2$  validated)
  - $\text{FIN}(T_1) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_1). C \cap A = \emptyset$ .
  - $\text{FIN}(T_1) > \text{VAL}(T_3). \text{WS}(T_3) \cap \text{WS}(T_1). B \cap A = \emptyset$ .
  - $\text{FIN}(T_2) > \text{START}(T_3). \text{RS}(T_3) \cap \text{WS}(T_2). (C) \cap (C) = C$ .  $T_3$  rolled back.