

## **Chap 19 solutions**

### **19.1.1**

The difference between strict and nonstrict locking for this example is only in the matter of whether the lock on  $A$  is released prior to the write of  $B$ . That is, the lock on  $A$  must be taken before  $r_1(A)$ . The lock on  $B$  can occur in any of the three positions prior to  $r_1(B)$  (3 choices).

If locking is strict, then the two unlocks must occur, in either order, after the second write action. Thus, there are 6 orders for strict, 2-phase locking.

If locking is not strict, the unlock of  $B$  must occur after  $w_1(B)$ , but then the unlock of  $A$  can occur in any of the three positions after  $w_1(A)$ . Thus, there are 9 orders if strictness is not required, and 3 of these must be nonstrict.

### **19.1.2a)**

$T_1$  wrote only  $B$ , but that value was later read by  $T_3$ . Thus,  $T_3$  must be rolled back.  $T_3$  wrote  $D$ , but no transaction has read  $D$ , so no further rollbacks are needed.

### **19.1.2b)**

$T_2$  rolled back since it reads  $B$  written by  $T_1$ .

$T_3$  rolled back since it reads  $C$  written by  $T_2$ .

### **19.1.2c)**

$T_2$  rolled back since it reads  $B$  written by  $T_1$ .

$T_3$  rolled back since it reads  $B$  written by  $T_1$ .

### **19.1.2d)**

$T_3$  rolled back as it reads  $B$  written by  $T_1$ .

### **19.1.3a)**

Inserting commit actions we get the following schedule:

$r_1(A)r_2(B)w_1(B)c_1w_2(C)c_2r_3(B)r_3(C)w_3(D)c_3$

- i) If the lost tail of log begins before  $c_1$ ,  $T_1$ ,  $T_2$  and  $T_3$  are considered uncommitted. If lost tail starts between  $c_1$  and  $c_2$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail starts between  $c_2$  and  $c_3$ ,  $T_3$  considered uncommitted.
- ii) No dirty reads
- iii) If  $c_1$  lost, dirty read of  $B$  by  $T_3$  since it is written by  $T_1$ .

### **19.1.3b)**

$r_1(A)w_1(B)c_1r_2(B)w_2(C)c_2r_3(C)w_3(D)c_3$

- i) If lost tail begins before  $c_1$ ,  $T_1$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail starts between  $c_1$  and  $c_2$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail between  $c_2$  and  $c_3$ ,  $t_3$  considered uncommitted.
- ii) No dirty reads
- iii) If  $c_1$  lost, dirty read of B by  $T_2$  since it is written by  $T_1$ . If  $c_2$  lost dirty read of C by  $T_3$  since it is written by  $T_2$ .

### **19.1.3c)**

$r_2(A)r_3(A)r_1(A)w_1(B)c_1r_2(B)r_3(B)w_2(C)c_2r_3(C)c_3$

- i) If lost tail begins before  $c_1$ ,  $T_1$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail begins between  $c_1$  and  $c_2$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail begins between  $c_2$  and  $c_3$ ,  $T_3$  considered uncommitted.
- ii) No dirty reads
- iii) If lost record is  $c_1$ , dirty read of B by  $T_2$  since it is written by  $T_1$ . If lost record if  $c_2$ , then dirty read of C by  $T_3$  since it is written by  $T_2$ .

### **19.1.3d)**

$r_2(A)r_3(A)r_1(A)w_1(B)c_1r_3(B)w_2(C)c_2r_3(C)c_3$

- i) If lost tail before  $c_1$ ,  $T_1$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail between  $c_1$  and  $c_2$ ,  $T_2$  and  $T_3$  considered uncommitted. If lost tail between  $c_2$  and  $c_3$ ,  $T_3$  considered uncommitted.
- ii) No dirty reads
- iii) If  $c_1$  lost, dirty read of B by  $T_2$  since it is written by  $T_1$ . If  $c_2$  lost, dirty read of C by  $T_3$  since it is written by  $T_2$ .

### **19.1.4a)**

For Recoverable schedules, each transaction must commit only after the transaction it read from commits. There is one read in each transaction.

Therefore the constraints are

- (i) When  $r_1(C)$  follows  $w_2(C)$ ,  $c_1$  must follow  $c_2$
- (ii) When  $r_2(B)$  follows  $w_1(B)$ ,  $c_2$  must follow  $c_1$

There are 10 transactions that violate (i) viz.

$w_2(A) r_2(B) w_2(C) w_1(A) w_1(B) r_1(C) c_1 c_2$   
 $w_2(A) r_2(B) w_1(A) w_2(C) w_1(B) r_1(C) c_1 c_2$   
 $w_2(A) r_2(B) w_1(A) w_1(B) w_2(C) r_1(C) c_1 c_2$   
 $w_2(A) w_1(A) r_2(B) w_2(C) w_1(B) r_1(C) c_1 c_2$   
 $w_2(A) w_1(A) r_2(B) w_1(B) w_2(C) r_1(C) c_1 c_2$   
 $w_2(A) w_1(A) w_1(B) r_2(B) w_2(C) r_1(C) c_1 c_2$   
 $w_1(A) w_2(A) r_2(B) w_2(C) w_1(B) r_1(C) c_1 c_2$   
 $w_1(A) w_2(A) r_2(B) w_1(B) w_2(C) r_1(C) c_1 c_2$   
 $w_1(A) w_2(A) w_1(B) r_2(B) w_2(C) r_1(C) c_1 c_2$   
 $w_1(A) w_1(B) w_2(A) r_2(B) w_2(C) r_1(C) c_1 c_2$

There are 13 transactions that violate (ii) viz.

w2(A) w1(A) w1(B) r2(B) w2(C) c\_2 r1(C) c\_1  
 w2(A) w1(A) w1(B) r2(B) w2(C) r1(C) c\_2 c\_1  
 w2(A) w1(A) w1(B) r2(B) r1(C) w2(C) c\_2 c\_1  
 w2(A) w1(A) w1(B) r1(C) r2(B) w2(C) c\_2 c\_1  
 w1(A) w2(A) w1(B) r2(B) w2(C) c\_2 r1(C) c\_1  
 w1(A) w2(A) w1(B) r2(B) w2(C) r1(C) c\_2 c\_1  
 w1(A) w2(A) w1(B) r2(B) r1(C) w2(C) c\_2 c\_1  
 w1(A) w2(A) w1(B) r1(C) r2(B) w2(C) c\_2 c\_1  
 w1(A) w1(B) w2(A) r2(B) w2(C) c\_2 r1(C) c\_1  
 w1(A) w1(B) w2(A) r2(B) w2(C) r1(C) c\_2 c\_1  
 w1(A) w1(B) w2(A) r2(B) r1(C) w2(C) c\_2 c\_1  
 w1(A) w1(B) w2(A) r1(C) r2(B) w2(C) c\_2 c\_1  
 w1(A) w1(B) r1(C) w2(A) r2(B) w2(C) c\_2 c\_1

Total possible orders are  $(8 \text{ choose } 4) = 70$ .

Out of these, 23 (13+10) are not recoverable.

Thus 47 (70 - 23) are recoverable.

#### **19.1.4b)**

In an ACR schedule, transactions only read values written by committed transactions.

Again each transaction has one read.

Therefore the constraints are

- (i) When r1(C) follows w2(c), r1(c) must also follow c2
- (ii) When r2(B) follows w1(B), r2(B) must also follow c1

18 schedules that violate (i) only are

w2(A) r2(B) w2(C) w1(A) w1(B) r1(C) c\_2 c\_1  
 w2(A) r2(B) w2(C) w1(A) w1(B) r1(C) c\_1 c\_2  
 w2(A) r2(B) w1(A) w2(C) w1(B) r1(C) c\_2 c\_1  
 w2(A) r2(B) w1(A) w2(C) w1(B) r1(C) c\_1 c\_2  
 w2(A) r2(B) w1(A) w1(B) w2(C) r1(C) c\_2 c\_1  
 w2(A) r2(B) w1(A) w1(B) w2(C) r1(C) c\_1 c\_2  
 w2(A) w1(A) r2(B) w2(C) w1(B) r1(C) c\_2 c\_1  
 w2(A) w1(A) r2(B) w2(C) w1(B) r1(C) c\_1 c\_2  
 w2(A) w1(A) r2(B) w1(B) w2(C) r1(C) c\_2 c\_1  
 w2(A) w1(A) r2(B) w1(B) w2(C) r1(C) c\_1 c\_2  
 w1(A) w2(A) r2(B) w2(C) w1(B) r1(C) c\_2 c\_1  
 w1(A) w2(A) r2(B) w2(C) w1(B) r1(C) c\_1 c\_2  
 w1(A) w2(A) r2(B) w1(B) w2(C) r1(C) c\_2 c\_1  
 w1(A) w2(A) r2(B) w1(B) w2(C) r1(C) c\_1 c\_2  
 w1(A) w2(A) w1(B) r2(B) w2(C) r1(C) c\_2 c\_1  
 w1(A) w2(A) w1(B) r2(B) w2(C) r1(C) c\_1 c\_2  
 w1(A) w1(B) w2(A) r2(B) w2(C) r1(C) c\_2 c\_1

w1(A) w1(B) w2(A) r2(B) w2(C) r1(C) c\_1 c\_2

9 schedules that violate only (ii) are

w2(A) w1(A) w1(B) r2(B) w2(C) c\_2 r1(C) c\_1  
w2(A) w1(A) w1(B) r2(B) r1(C) w2(C) c\_2 c\_1  
w2(A) w1(A) w1(B) r2(B) r1(C) w2(C) c\_1 c\_2  
w2(A) w1(A) w1(B) r2(B) r1(C) c\_1 w2(C) c\_2  
w2(A) w1(A) w1(B) r1(C) r2(B) w2(C) c\_2 c\_1  
w2(A) w1(A) w1(B) r1(C) r2(B) w2(C) c\_1 c\_2  
w2(A) w1(A) w1(B) r1(C) r2(B) c\_1 w2(C) c\_2  
w1(A) w2(A) w1(B) r2(B) w2(C) c\_2 r1(C) c\_1  
w1(A) w1(B) w2(A) r2(B) w2(C) c\_2 r1(C) c\_1

Also schedules that cause violations of both (i) and (ii) are

w2(A) w1(A) w1(B) r2(B) w2(C) r1(C) c\_2 c\_1  
w2(A) w1(A) w1(B) r2(B) w2(C) r1(C) c\_1 c\_2

Thus there are 41 ACR schedules (70 - 18 - 9 - 2).

#### **19.1.4c)**

All corresponding actions of the two transactions on same element are possible conflicts  
i.e. cannot be swapped without affecting the serializable nature.

Of the 47 recoverable schedules, below 19 are not serializable due to conflicts.

w2(A) r2(B) w1(A) w1(B) r1(C) w2(C) c\_2 c\_1  
w2(A) r2(B) w1(A) w1(B) r1(C) w2(C) c\_1 c\_2  
w2(A) r2(B) w1(A) w1(B) r1(C) c\_1 w2(C) c\_2  
w2(A) w1(A) r2(B) w1(B) r1(C) w2(C) c\_2 c\_1  
w2(A) w1(A) r2(B) w1(B) r1(C) w2(C) c\_1 c\_2  
w2(A) w1(A) r2(B) w1(B) r1(C) c\_1 w2(C) c\_2  
w2(A) w1(A) w1(B) r2(B) r1(C) w2(C) c\_1 c\_2 .  
w2(A) w1(A) w1(B) r2(B) r1(C) c\_1 w2(C) c\_2 .  
w2(A) w1(A) w1(B) r1(C) r2(B) w2(C) c\_1 c\_2 .  
w2(A) w1(A) w1(B) r1(C) r2(B) c\_1 w2(C) c\_2 .  
w2(A) w1(A) w1(B) r1(C) c\_1 r2(B) w2(C) c\_2  
w1(A) w2(A) r2(B) w2(C) c\_2 w1(B) r1(C) c\_1  
w1(A) w2(A) r2(B) w2(C) w1(B) c\_2 r1(C) c\_1  
w1(A) w2(A) r2(B) w2(C) w1(B) r1(C) c\_2 c\_1 .  
w1(A) w2(A) r2(B) w1(B) w2(C) c\_2 r1(C) c\_1  
w1(A) w2(A) r2(B) w1(B) w2(C) r1(C) c\_2 c\_1 .  
w1(A) w2(A) r2(B) w1(B) r1(C) w2(C) c\_2 c\_1  
w1(A) w2(A) r2(B) w1(B) r1(C) w2(C) c\_1 c\_2  
w1(A) w2(A) r2(B) w1(B) r1(C) c\_1 w2(C) c\_2

Thus 28 transactions (47 - 19) are both recoverable and serializable.

#### **19.1.4d)**

Of the 41 ACR schedules, following 13 are not conflict serializable.

w2(A) r2(B) w1(A) w1(B) r1(C) w2(C) c<sub>2</sub> c<sub>1</sub>  
 w2(A) r2(B) w1(A) w1(B) r1(C) w2(C) c<sub>1</sub> c<sub>2</sub>  
 w2(A) r2(B) w1(A) w1(B) r1(C) c<sub>1</sub> w2(C) c<sub>2</sub>  
 w2(A) w1(A) r2(B) w1(B) r1(C) w2(C) c<sub>2</sub> c<sub>1</sub>  
 w2(A) w1(A) r2(B) w1(B) r1(C) w2(C) c<sub>1</sub> c<sub>2</sub>  
 w2(A) w1(A) r2(B) w1(B) r1(C) c<sub>1</sub> w2(C) c<sub>2</sub>  
 w2(A) w1(A) w1(B) r1(C) c<sub>1</sub> r2(B) w2(C) c<sub>2</sub>  
 w1(A) w2(A) r2(B) w2(C) c<sub>2</sub> w1(B) r1(C) c<sub>1</sub>  
 w1(A) w2(A) r2(B) w2(C) w1(B) c<sub>2</sub> r1(C) c<sub>1</sub>  
 w1(A) w2(A) r2(B) w1(B) w2(C) c<sub>2</sub> r1(C) c<sub>1</sub>  
 w1(A) w2(A) r2(B) w1(B) r1(C) w2(C) c<sub>2</sub> c<sub>1</sub>  
 w1(A) w2(A) r2(B) w1(B) r1(C) w2(C) c<sub>1</sub> c<sub>2</sub>  
 w1(A) w2(A) r2(B) w1(B) r1(C) c<sub>1</sub> w2(C) c<sub>2</sub>

Thus 28 transactions (41 - 13) are both ACR and serializable.

#### **19.1.5**

Example of ACR schedule (tx reads values only written by committed trans):

w<sub>1</sub>(A)w<sub>1</sub>(B)w<sub>2</sub>(A)c<sub>1</sub>r<sub>2</sub>(B)c<sub>2</sub>

We insert shared and exclusive lock such that the schedule becomes non-strict(releases exclusive lock before commit/abort)

xl<sub>1</sub>(A)w<sub>1</sub>(A)ul<sub>1</sub>(A)xl<sub>1</sub>(B)w<sub>1</sub>(B)ul<sub>1</sub>(B)xl<sub>2</sub>(A)w<sub>2</sub>(A) ul<sub>2</sub>(A)c<sub>1</sub>sl<sub>2</sub>(B)r<sub>2</sub>(B)ul<sub>2</sub>(B)c<sub>2</sub>


We see that T<sub>1</sub> release X lock on A (ul<sub>1</sub>(A) before c<sub>1</sub>.

#### **19.2.1a)**


After inserting the appropriate lock and unlock requests, the schedule becomes(assuming no delays):

sl<sub>1</sub>(A)r<sub>1</sub>(A)sl<sub>2</sub>(B)r<sub>2</sub>(B)xl<sub>1</sub>(C)w<sub>1</sub>(C)sl<sub>3</sub>(D)r<sub>3</sub>(D)sl<sub>4</sub>(E)r<sub>4</sub>(E)xl<sub>3</sub>(B)w<sub>3</sub>(B)ul<sub>3</sub>(B,D)xl<sub>2</sub>(C)w<sub>2</sub>(C)ul<sub>2</sub>(C,B)xl<sub>4</sub>(A)w<sub>4</sub>(A)ul<sub>4</sub>(A,E)xl<sub>1</sub>(D)w<sub>1</sub>(D)ul<sub>1</sub>(D,A,C)


- sl<sub>1</sub>(A)r<sub>1</sub>(A)sl<sub>2</sub>(B)r<sub>2</sub>(B)xl<sub>1</sub>(C)w<sub>1</sub>(C) sl<sub>3</sub>(D)r<sub>3</sub>(D)sl<sub>4</sub>(E)r<sub>4</sub>(E) proceed fine.
- xl<sub>3</sub>(B) is denied since T<sub>2</sub> is holding a lock to B.




- xl<sub>2</sub>(C) is denied since T<sub>1</sub> is holding a lock to C



- $xl_4(A)$  is denied since  $T_1$  is holding a lock to A



- $xl_1(D)$  is denied since  $T_3$  is holding a lock to D. This causes a deadlock (see the cycle in the waits for graph).




- Assume  $T_1$  is aborted. Releases locks on A and C such that  $T_2$  and  $T_4$  can continue.


### **19.2.1b)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)xl_1(B)w_1(B)ul_1(B,A)xl_2(C)w_2(C)ul_2(C,B)xl_3(D)w_3(D)ul_3(D,C)$

- $sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)$  proceed fine
- $xl_1(B)$  is denied since  $T_2$  is holding lock to B



- $xl_2(C)$  is denied since  $T_3$  is holding a lock on C.




- $xl_3(D)w_3(D)ul_3(D,C)$  complete.
- After  $T_3$  finishes,  $T_2$  completes followed by  $T_1$ .


### **19.2.1c)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)xl_1(B)w_1(B)ul_1(A,B)xl_2(C)w_2(C)ul_2(C,B)xl_3(A)w_3(A)ul_3(A,C)$


- $sl_1(A)r_1(A)sl_2(B)r_2(B)sl_3(C)r_3(C)$  proceed fine.
- $xl_1(B)$  is denied since  $T_2$  is holding a lock on B



- $xl_2(C)$  is denied since  $T_3$  holds a lock on C



- $xl_3(A)$  is denied since  $T_1$  is holding a lock on A. There is a deadlock as shown by the cycle in the waits-for graph.




- If  $T_3$  is aborted, releases lock on C.  $T_2$  can continue. After  $T_2$ ,  $T_1$  can continue.


### **19.2.1d)**

$sl_1(A)r_1(A)sl_2(B)r_2(B)xl_1(C)w_1(C)xl_2(D)w_2(D)sl_3(C)r_3(C)ul_3(C)xl_1(B)w_1(B)ul_1(B,A,C)xl_4(D)ul_4(D)xl_2(A)w_2(A)ul_2(A,B,D)$


- $sl_1(A)r_1(A)sl_2(B)r_2(B)xl_1(C)w_1(C)xl_2(D)w_2(D)$  proceed fine.
- $sl_3(C)$  is denied since  $T_1$  is holding a lock on C.




- $xl_1(B)$  is denied since  $T_2$  is holding a lock on B.



- $xl_4(D)$  is denied since  $T_2$  is holding a lock on D.



- $xl_2(A)$  is denied since  $T_1$  is holding a lock on A. This creates a deadlock as shown by the cycle.



- Assume  $T_2$  is aborted. Then  $T_1$  and  $T_4$  can continue. After  $T_1$ ,  $T_3$  can continue.

#### 19.2.2a)

<u><b>T<sub>1</sub></b></u>	<u><b>T<sub>2</sub></b></u>	<u><b>T<sub>3</sub></b></u>	<u><b>T<sub>4</sub></b></u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)			
	sl <sub>2</sub> (B)r <sub>2</sub> (B)		
xl <sub>1</sub> (C)w <sub>1</sub> (C)			
		sl <sub>3</sub> (D)r <sub>3</sub> (D)	
			sl <sub>4</sub> (E)r <sub>4</sub> (E)
		xl <sub>3</sub> (B) <b>Waits</b> since $T_2$ is older	
	xl <sub>2</sub> (C). <b>Waits</b> since $T_1$ is older.		
			xl <sub>4</sub> (A). <b>Waits</b> since $T_1$ is older.
xl <sub>1</sub> (D). Aborts $T_3$ since $T_1$ is older.		<b>Wounded.</b>	

- After  $T_1$  completes,  $T_2$  and  $T_4$  continue. Then  $T_3$  restarts and completes without interference.

#### 19.2.2b)

<u><b>T<sub>1</sub></b></u>	<u><b>T<sub>2</sub></b></u>	<u><b>T<sub>3</sub></b></u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)		
	sl <sub>1</sub> (B)r <sub>2</sub> (B)	
		sl <sub>3</sub> (C)r <sub>3</sub> (C)
xl <sub>1</sub> (B). Aborts $T_2$ since $T_1$ is older than $T_2$ . Completes.	<b>Wounded.</b>	
		xl <sub>3</sub> (D)w <sub>3</sub> (D). Completes.
	$T_2$ restarts and completes.	

#### 19.2.2c)

<u><b>T<sub>1</sub></b></u>	<u><b>T<sub>2</sub></b></u>	<u><b>T<sub>3</sub></b></u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)		
	sl <sub>2</sub> (B)r <sub>2</sub> (B)	
		sl <sub>3</sub> (C)r <sub>3</sub> (C)
xl <sub>1</sub> (B). Aborts $T_2$ since $T_1$ older than $T_2$ . w <sub>1</sub> (B).	<b>Wounded.</b>	



Completes.		
		xl <sub>3</sub> (A)w <sub>3</sub> (A). Completes.
	T <sub>2</sub> restarts and completes.	

#### 19.2.2d)

<u>T1</u>	<u>T2</u>	<u>T3</u>	<u>T4</u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)			
	sl <sub>2</sub> (B)r <sub>2</sub> (B)		
xl <sub>1</sub> (C)w <sub>1</sub> (C)			
	xl <sub>2</sub> (D)w <sub>2</sub> (D)		
		sl <sub>3</sub> (C)r <sub>3</sub> (C). <b>Waits</b> since T <sub>1</sub> is older.	
xl <sub>1</sub> (B). Aborts T <sub>2</sub> since T <sub>1</sub> is older. w <sub>1</sub> (B). Completes.	<b>Wounded</b>		
		sl <sub>3</sub> (C)r <sub>3</sub> (C). Completes.	
			xl <sub>4</sub> (D). Completes.
	T <sub>2</sub> restarts and completes.		

#### 19.2.3a)

<u>T1</u>	<u>T2</u>	<u>T3</u>	<u>T4</u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)			
	sl <sub>2</sub> (B)r <sub>2</sub> (B)		
xl <sub>1</sub> (C)w <sub>1</sub> (C)			
		sl <sub>3</sub> (D)r <sub>3</sub> (D)	
			sl <sub>4</sub> (E)r <sub>4</sub> (E)
		xl <sub>3</sub> (B). <b>Dies</b> since T <sub>2</sub> is older.	
	xl <sub>2</sub> (C). <b>Dies</b> since T <sub>1</sub> is older.		
			xl <sub>4</sub> (A). <b>Dies</b> since T <sub>1</sub> is older.
xl <sub>1</sub> (D).w <sub>1</sub> (D). Completes.			
	T <sub>2</sub> restarts. sl <sub>2</sub> (B)r <sub>2</sub> (B)		
		T <sub>3</sub> restarts. sl <sub>3</sub> (D)r <sub>3</sub> (D)	
			T <sub>4</sub> restarts. sl <sub>4</sub> (E)r <sub>4</sub> (E)
	xl <sub>2</sub> (C)w <sub>2</sub> (C). Completes		

		xl <sub>3</sub> (B)w <sub>3</sub> (B). Completes	
			xl <sub>4</sub> (A)w <sub>4</sub> (A). Completes.

### 19.2.3b)

<u>T1</u>	<u>T2</u>	<u>T3</u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)		
	sl <sub>2</sub> (B)r <sub>2</sub> (B)	
		sl <sub>3</sub> (C)r <sub>3</sub> (C)
xl <sub>1</sub> (B). Waits since T <sub>1</sub> is older than T <sub>2</sub> .		
	xl <sub>2</sub> (C). Waits since T <sub>2</sub> is older than T <sub>3</sub> .	
		xl <sub>3</sub> (D)w <sub>3</sub> (D). Completes.
	T <sub>2</sub> resumes and completes.	
T <sub>1</sub> resumes and completes.		

### 19.2.3c)

<u>T1</u>	<u>T2</u>	<u>T3</u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)		
	sl <sub>2</sub> (B)r <sub>2</sub> (B)	
		sl <sub>3</sub> (C)r <sub>3</sub> (C)
xl <sub>1</sub> (B). <b>Waits</b> since T <sub>1</sub> is older than T <sub>2</sub> .		
	xl <sub>2</sub> (C). <b>Waits</b> since T <sub>2</sub> is older than T <sub>3</sub> .	
		xl <sub>3</sub> (A). <b>Dies</b> since T <sub>1</sub> is older.
	T <sub>2</sub> resumes and completes.	
T <sub>1</sub> resumes and completes.		
		T <sub>3</sub> restarts and finishes.

### 19.2.3d)

<u>T1</u>	<u>T2</u>	<u>T3</u>	<u>T4</u>
sl <sub>1</sub> (A)r <sub>1</sub> (A)			
	sl <sub>2</sub> (B)r <sub>2</sub> (B)		
xl <sub>1</sub> (C)w <sub>1</sub> (C)			
	xl <sub>2</sub> (D)w <sub>2</sub> (D)		
		sl <sub>3</sub> (C). <b>Dies</b> since T <sub>1</sub> is older.	
xl <sub>1</sub> (B). <b>Waits</b> since T <sub>1</sub> is older			

than T <sub>2</sub> .			
			xl <sub>4</sub> (D). <b>Dies</b> since T <sub>2</sub> is older.
	xl <sub>2</sub> (A). <b>Dies</b> since T <sub>1</sub> is older.		
T <sub>1</sub> completes after T <sub>2</sub> dies.			
	T <sub>2</sub> restarts. sl <sub>2</sub> (B)r <sub>2</sub> (B)		
		T <sub>3</sub> restarts. sl <sub>3</sub> (C)r <sub>3</sub> (C). Completes.	
			T <sub>4</sub> restarts. xl <sub>4</sub> (D)w <sub>4</sub> (D). Completes.
	xl <sub>2</sub> (D)w <sub>2</sub> (D)xl <sub>2</sub> (A)w <sub>2</sub> (A). Completes		

Given T<sub>1</sub>, T<sub>2</sub>, ..., T<sub>n</sub>, for n > 1,

#### 19.2.4

There exists a {T<sub>1</sub>, T<sub>2</sub>, ..., T<sub>n</sub>} of waiting transaction such as T<sub>1</sub> is waiting for an item held by T<sub>2</sub>, T<sub>2</sub> is waiting for an item held by T<sub>3</sub>, ..., T<sub>n</sub> is waiting for an item held by T<sub>1</sub>. Thus, it is possible to have a wait-for graph with a cycle of length n, for n > 1. However, for n=1, there is no cycle is formed for a node since a transaction doesn't have to wait for a lock held by itself.


#### 19.2.5

Yes this approach would avoid deadlocks since a transaction would never start until all locks become available. So there would not be a situation where 2 transactions are waiting for each other for locking a resources. Timeouts can be added to prevent starvation.

#### 19.2.6

To construct a waits-for graph for an intention-locking system, the arcs need to be drawn such that:

- the arc goes from transaction waiting for a lock on a node to the transaction holding the lock on the node or its child node that causes the first transaction to wait. For eg., In the tree structure shown, if T<sub>1</sub> holds an IX lock on A because it holds a X lock on B (to updated B), and then T<sub>2</sub> comes in with a write of A (requiring a X on A), then the arc is drawn from T<sub>2</sub> to T<sub>1</sub>. If multiple transactions hold IX lock on A (for eg, due to write of C by T<sub>3</sub>), an arc would be drawn from T<sub>2</sub> to T<sub>3</sub> as well. Similar holds true if T<sub>1</sub> and T<sub>3</sub> were holding IS locks on A, and T<sub>3</sub> came in with a request for X lock on A. Basically the arcs needs to be drawn from the transaction that has to wait for a lock on a node to all the transaction that are holding locks on that node that causes the incoming transaction to waits.



### 19.2.7)

Suppose  $T_1$  is a transaction that tries to lock elements  $A$  and  $B$  in that order. There are also an indefinitely long sequence of transactions  $T_2, T_3, \dots$  that lock  $B$  and then  $A$ .

Initially,  $T_1$  locks  $A$  and  $T_2$  locks  $B$  and then requests a lock on  $A$ .  $T_2$  therefore waits for  $T_1$ . Meanwhile,  $T_3$  starts and requests a lock on  $B$ . Thus,  $T_3$  waits for  $T_2$ , and the waits-for graph is:

$T_3 \text{ ----> } T_2 \text{ ----> } T_1$

When  $T_1$  requests a lock on  $A$ , it would cause a cycle, so  $T_1$  is rolled back. Now,  $T_2$  completes and releases its locks. The lock on  $A$  is given to  $T_1$  and the lock on  $B$  is given to  $T_3$ . Next,  $T_3$  requests a lock on  $A$ , but has to wait for  $T_1$ . Also,  $T_4$  starts up and requests a lock on  $B$ , thus being forced to wait for  $T_3$ . The waits-for graph is now:

$T_4 \text{ ----> } T_3 \text{ ----> } T_1$

The above graph is just like the first, except  $T_4$  and  $T_3$  have replaced  $T_3$  and  $T_2$ , respectively. As long as there is a supply of transactions like  $T_2, T_3$ , and  $T_4$ , they can prevent  $T_1$  from finishing ever.

Notice that the transactions in the example above request locks on  $A$  and  $B$  in different orders. Perhaps forcing transactions to request locks in a fixed order will solve the problem. Indeed, if we follow the policy that when a lock becomes available, it is given to a waiting transaction on a first-come-first-served basis, then any transaction that is waiting for a lock will eventually become the transaction that has been waiting for the lock the longest. At that point, the next time the lock becomes available, it will be granted to that transaction, which thus makes some progress. The above observation is the germ of an inductive proof that requesting locks in order plus first-come-first-served will allow every transaction eventually to complete.

However, if the scheduler is able to give locks to any transaction it wishes, then it is easy to make a transaction starve. Suppose  $T_1$  wants a lock on  $A$  when some other transaction has that lock. As long as there is always another transaction besides  $T_1$  waiting for a lock on  $A$ , and the scheduler always chooses to give the lock to some transaction other than  $T_1$  when the lock becomes available,  $T_1$  never makes any progress.

Last, let us consider deadlock prevention by timeout. If the duration before a timeout can vary even slightly, then the example given initially for deadlock prevention by avoiding cycles will work. That is, even though  $T_2$  has been waiting slightly longer for a lock than

$T_1$ , if we allow that  $T_1$  might timeout first, then the sequence of events described above could also occur with timeout-based deadlock prevention.

If timeouts occur for any transaction at exactly  $t$  seconds from when it first started to wait, then we need a simple modification of the above example. Assume that  $T_1$  is faster at requesting its second lock than any of the other transactions. Then, after  $T_1$  gets a lock on  $A$  and  $T_2$  gets a lock on  $B$ ,  $T_1$  next requests its lock on  $B$  and starts to wait. Slightly later,  $T_2$  requests a lock on  $A$ .  $T_3$  requests its lock on  $B$  just before  $T_1$  times out, so when that timeout occurs,  $T_2$  completes, gives its  $A$ -lock to  $T_1$  and its  $B$ -lock to  $T_3$ . When  $T_2$  next requests a lock on  $B$  and starts to wait, this cycle can repeat with  $T_3$  and a new  $T_4$  in place of  $T_2$  and  $T_3$ .

### **19.3.5**

The task of the compensating transaction is first to determine whether the file  $f$  is still the present one with that name, or whether it has been replaced with a later version.

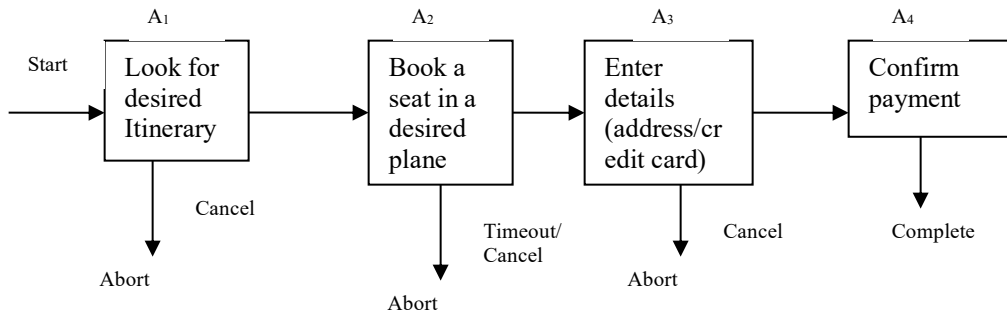
1. If  $f$  is still in place, then the compensating transaction must
  1. Restore  $f'$  if the latter existed, or
  2. Delete  $f$  if not.
2. If  $f$  has been overwritten, then there is no need to change the file with that name.

To see that this compensation works, we can consider all the possible cases above, and check that the file system, after compensation, is the same as would be the case had  $f$  never been written. If  $f$  was overwritten before compensation, then we surely leave the file system in the same state as if  $f$  were never written. If  $f$  is not overwritten, then we have correctly left the file with that name being either  $f'$ , if it existed, or left the file system with no file of that name.

There are two interesting issues. First, in order for there to be a compensating transaction, it seems we must both remember the time at which the installation occurred (so we can tell whether the file  $f$  has been overwritten when the uninstallation occurs), and we must also remember the file  $f'$ , at least until such time as the file is overwritten, and the new value is itself guaranteed never to be uninstalled (i.e., the saga that wrote it has finished). These features are not commonly available.

Second, note that file systems do not support serializability. Thus, it is entirely possible that the uninstalled file  $f$  has been read, and its contents have influenced other files or the behavior of the system.

### **19.3.2**



List of compensating actions:

- A1<sup>-</sup> : Remove desired itinerary details from database (if saved)
- A2<sup>-</sup>: Put booked seat back in the available pool of seats
- A3<sup>-</sup>: Remove details (address/credit card) from database (if saved)
- A4<sup>-</sup>: Refund charged amount to credit card